

Lab2: Linear Regression

Mehreen Ali Gillani

2026-02-11

```
# Import necessary libraries
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from scipy import stats
# Load the dataset
url = 'https://raw.githubusercontent.com/georgehagstrom/DATA622Spring2026/main/website/assign
earnings = pd.read_csv(url)
# Display the first few rows of the dataset
earnings.head()
```

	height	weight	male	earn	earnk	ethnicity	education	mother_education	father_education
0	74	210.0	1	50000.0	50.0	White	16.0	12.0	12.0
1	66	125.0	0	60000.0	60.0	White	16.0	12.0	12.0
2	64	126.0	0	30000.0	30.0	White	16.0	12.0	12.0
3	65	200.0	0	25000.0	25.0	White	17.0	12.0	NaN
4	63	110.0	0	50000.0	50.0	Other	16.0	15.0	12.0

- a) EDA: Perform a basic exploratory data analysis on this dataset. Tell me how you decided to handle rows with missing values, and show me the plots of the distributions of variables and the relationships between variables that you believe are most important. An important decision is how to handle earnk, the target variable, and whether to log transform it or not. Look at the distribution of earnk and discuss the pros and cons of using the log transform, and make a decision about whether perform to the transformation or not.

EDA

Missing Values Analysis

```
# Check for missing values and their percentages
print((earnings.isnull().mean() * 100).round(2))
```

```
height          0.00
weight         1.49
male           0.00
earn           0.00
earnk          0.00
ethnicity      0.00
education      0.11
mother_education 13.44
father_education 16.24
walk           0.00
exercise        0.00
smokenow       0.06
tense          0.11
angry          0.00
age            0.00
dtype: float64
```

Minimal missing (<1%): education (2), smokenow (1), tense (1), angry (1) Moderate missing: weight (27), 1.49% Substantial missing: mother_education (244), father_education (295) 13.44% and 16.24% respectively # Decision: Drop rows with missing values in the target variable ‘earnk’ and moderate missing ‘weight’, but impute or drop the substantial missing ‘mother_education’ and ‘father_education’ # Drop rows with missing values in ‘earnk’ and ‘weight’

```
# import necessary libraries for imputation
from sklearn.impute import SimpleImputer
from sklearn.experimental import enable_iterative_imputer
from sklearn.impute import IterativeImputer
import statsmodels.api as sm

earnings_cleaned = earnings.copy()
# Mode imputation for categorical variables with minimal missingness
earnings_cleaned['angry'].fillna(earnings_cleaned['angry'].mode()[0], inplace=True)
```

```

earnings_cleaned['tense'].fillna(earnings_cleaned['tense'].mode()[0], inplace=True)

# Mode imputation for 'education' and 'smokenow' with minimal missingness
imputer = SimpleImputer(strategy='most_frequent')
earnings_cleaned[['education', 'smokenow']] = imputer.fit_transform(earnings_cleaned[['educa

# Create imputation model using complete variables
imputer = IterativeImputer(max_iter=10, random_state=0)
earnings_cleaned['weight'] = imputer.fit_transform(earnings_cleaned[['weight', 'height', 'ma
print((earnings_cleaned.isnull().mean() * 100).round(2))

# Create missing indicators and impute in-place
earnings_cleaned['mother_edu_missing'] = earnings_cleaned['mother_education'].isnull().astype
earnings_cleaned['father_edu_missing'] = earnings_cleaned['father_education'].isnull().astype
earnings_cleaned['mother_education'] = earnings_cleaned['mother_education'].fillna(earnings_
earnings_cleaned['father_education'] = earnings_cleaned['father_education'].fillna(earnings_

# Verify
print(f"\nVerification - any remaining missing?")
print(f"Mother education: {earnings_cleaned['mother_education'].isnull().sum()}")
print(f"Father education: {earnings_cleaned['father_education'].isnull().sum()}")

```

height	0.00
weight	0.00
male	0.00
earn	0.00
earnk	0.00
ethnicity	0.00
education	0.00
mother_education	13.44
father_education	16.24
walk	0.00
exercise	0.00
smokenow	0.00
tense	0.00
angry	0.00
age	0.00
	dtype: float64

Verification - any remaining missing?

```
Mother education: 0
```

```
Father education: 0
```

```
/var/folders/6w/rkqvbjtx04s93768hcryddtc0000gn/T/ipykernel_63361/1825402175.py:9: FutureWarning:
```

```
The behavior will change in pandas 3.0. This inplace method will never work because the inter-
```

```
For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: va-
```

```
earnings_cleaned['angry'].fillna(earnings_cleaned['angry'].mode()[0], inplace=True)
```

```
/var/folders/6w/rkqvbjtx04s93768hcryddtc0000gn/T/ipykernel_63361/1825402175.py:10: FutureWarning:
```

```
The behavior will change in pandas 3.0. This inplace method will never work because the inter-
```

```
For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: va-
```

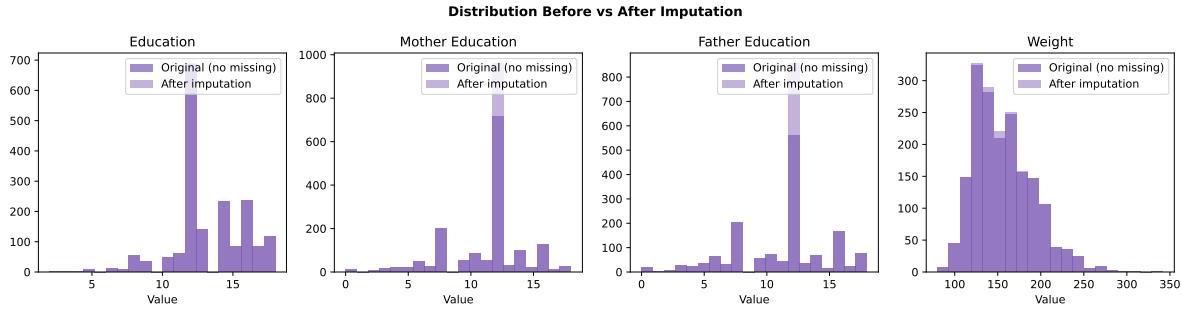
```
earnings_cleaned['tense'].fillna(earnings_cleaned['tense'].mode()[0], inplace=True)
```

Before and After Imputation Distributions

```
fig, axes = plt.subplots(1, 4, figsize=(15, 4))

for i, var in enumerate(['education', 'mother_education', 'father_education', 'weight']):
    # Plot before (original with missing removed) vs after (imputed)
    axes[i].hist(earnings[var].dropna(), bins=20, alpha=0.5, label='Original (no missing)', color='blue')
    axes[i].hist(earnings_cleaned[var], bins=20, alpha=0.5, label='After imputation', color='red')
    axes[i].set_title(f'{var.replace("_", " ").title()}'')
    axes[i].set_xlabel('Value')
    axes[i].legend()

plt.suptitle('Distribution Before vs After Imputation', fontweight='bold')
plt.tight_layout()
plt.show()
```



Summary statistics before and after imputation

```
# Before and after imputation summary

for var in ['mother_education', 'father_education', 'weight', 'education']:
    if earnings[var].isnull().sum() > 0:
        print(f'{var:20} Standard deviation:{earnings[var].std():6.1f} → {earnings_cleaned[var].std():6.1f}')
        print(f'Mean:{earnings[var].mean():6.1f} → {earnings_cleaned[var].mean():6.1f} ')
        print(f'Median:{earnings[var].median():6.1f} → {earnings_cleaned[var].median():6.1f}')


mother_education      Standard deviation:   3.2 →   2.9 Mean:  11.2 →  11.3 Median: 12.0 → 12.0
father_education       Standard deviation:   3.7 →   3.4 Mean:  11.2 →  11.3 Median: 12.0 → 12.0
weight                  Standard deviation: 34.6 → 34.4 Mean: 156.3 → 156.1 Median: 150.0 → 150.0
education                Standard deviation:   2.6 →   2.6 Mean:  13.2 →  13.2 Median: 12.0 → 12.0
```

Visualizations

```
# Set style
sns.set_style("whitegrid")
plt.rcParams['figure.figsize'] = (15, 12)

# 1. DISTRIBUTION PLOTS - Key Variables
fig, axes = plt.subplots(3, 3, figsize=(15, 12))

# Earnings (target variable)
axes[0,0].hist(earnings_cleaned['earnk'], bins=30, edgecolor='black', alpha=0.7)
axes[0,0].set_title('Earnings Distribution (earnk)', fontsize=12, fontweight='bold')
axes[0,0].set_xlabel('Annual Earnings ($)')
axes[0,0].set_ylabel('Frequency')
```

```

# Education
axes[0,1].hist(earnings_cleaned['education'], bins=20, edgecolor='black', alpha=0.7, color='green')
axes[0,1].set_title('Education Distribution', fontsize=12, fontweight='bold')
axes[0,1].set_xlabel('Years of Education')
axes[0,1].set_ylabel('Frequency')

# Age
axes[0,2].hist(earnings_cleaned['age'], bins=30, edgecolor='black', alpha=0.7, color='orange')
axes[0,2].set_title('Age Distribution', fontsize=12, fontweight='bold')
axes[0,2].set_xlabel('Age')
axes[0,2].set_ylabel('Frequency')

# Parent Education (after imputation)
axes[1,0].hist(earnings_cleaned['mother_education'], bins=20, edgecolor='black', alpha=0.5, color='blue')
axes[1,0].hist(earnings_cleaned['father_education'], bins=20, edgecolor='black', alpha=0.5, color='red')
axes[1,0].set_title('Parent Education Distribution', fontsize=12, fontweight='bold')
axes[1,0].set_xlabel('Years of Education')
axes[1,0].set_ylabel('Frequency')
axes[1,0].legend()

# Height vs Weight
axes[1,1].scatter(earnings_cleaned['height'], earnings_cleaned['weight'], alpha=0.5, s=10)
axes[1,1].set_title('Height vs Weight', fontsize=12, fontweight='bold')
axes[1,1].set_xlabel('Height')
axes[1,1].set_ylabel('Weight')

# Gender distribution
gender_counts = earnings_cleaned['male'].value_counts()
axes[1,2].bar(['Female', 'Male'], gender_counts.values, color=['pink', 'lightblue'], edgecolor='black')
axes[1,2].set_title('Gender Distribution', fontsize=12, fontweight='bold')
axes[1,2].set_ylabel('Count')

# Missing indicators
missing_data = earnings_cleaned[['mother_edu_missing', 'father_edu_missing']].sum()
axes[2,0].bar(['Mother Edu Missing', 'Father Edu Missing'], missing_data.values, color=['red', 'blue'])
axes[2,0].set_title('Missing Data Indicators', fontsize=12, fontweight='bold')
axes[2,0].set_ylabel('Count')

# Exercise habits
if 'exercise' in earnings_cleaned.columns:
    exercise_counts = earnings_cleaned['exercise'].value_counts().sort_index()
    axes[2,1].bar(range(len(exercise_counts)), exercise_counts.values, edgecolor='black', color='purple')

```

```

axes[2,1].set_title('Exercise Frequency', fontsize=12, fontweight='bold')
axes[2,1].set_xlabel('Exercise Level')
axes[2,1].set_ylabel('Count')
axes[2,1].set_xticks(range(len(exercise_counts)))
axes[2,1].set_xticklabels(exercise_counts.index)

# Ethnicity distribution (top 5 only)
if 'ethnicity' in earnings_cleaned.columns:
    ethnicity_counts = earnings_cleaned['ethnicity'].value_counts().head(5)
    axes[2,2].bar(range(len(ethnicity_counts)), ethnicity_counts.values, edgecolor='black', color='lightblue')
    axes[2,2].set_title('Top 5 Ethnicities', fontsize=12, fontweight='bold')
    axes[2,2].set_ylabel('Count')
    axes[2,2].set_xticks(range(len(ethnicity_counts)))
    axes[2,2].set_xticklabels(ethnicity_counts.index, rotation=45, ha='right')

plt.tight_layout()
plt.suptitle('Key Variable Distributions', fontsize=16, fontweight='bold', y=1.02)
plt.show()

# 2. RELATIONSHIP PLOTS - Most Important Correlations
fig, axes = plt.subplots(2, 2, figsize=(14, 10))

# Plot 1: Earnings vs Education (THE MOST IMPORTANT)
axes[0,0].scatter(earnings_cleaned['education'], earnings_cleaned['earnk'], alpha=0.4, s=20)
# Add trend line
z = np.polyfit(earnings_cleaned['education'], earnings_cleaned['earnk'], 1)
p = np.poly1d(z)
education_range = np.linspace(earnings_cleaned['education'].min(), earnings_cleaned['education'].max(), 100)
axes[0,0].plot(education_range, p(education_range), "r--", linewidth=2, label=f'Trend: y={z[0]:.0f}x{z[1]:.0f}')
axes[0,0].set_title('Earnings vs Education', fontsize=14, fontweight='bold')
axes[0,0].set_xlabel('Years of Education')
axes[0,0].set_ylabel('Annual Earnings ($)')
axes[0,0].legend()
axes[0,0].grid(True, alpha=0.3)

# Plot 2: Earnings vs Age
axes[0,1].scatter(earnings_cleaned['age'], earnings_cleaned['earnk'], alpha=0.4, s=20, color='lightblue')
# Add trend line
z = np.polyfit(earnings_cleaned['age'], earnings_cleaned['earnk'], 1)
p = np.poly1d(z)
age_range = np.linspace(earnings_cleaned['age'].min(), earnings_cleaned['age'].max(), 100)
axes[0,1].plot(age_range, p(age_range), "r--", linewidth=2, label=f'Trend: y={z[0]:.0f}x{z[1]:.0f}')

```

```

axes[0,1].set_title('Earnings vs Age', fontsize=14, fontweight='bold')
axes[0,1].set_xlabel('Age')
axes[0,1].set_ylabel('Annual Earnings ($)')
axes[0,1].legend()
axes[0,1].grid(True, alpha=0.3)

# Plot 3: Earnings by Gender (Boxplot)
male_earnings = earnings_cleaned[earnings_cleaned['male'] == 1]['earnk']
female_earnings = earnings_cleaned[earnings_cleaned['male'] == 0]['earnk']
box_data = [female_earnings.dropna(), male_earnings.dropna()]
axes[1,0].boxplot(box_data, labels=['Female', 'Male'])
axes[1,0].set_title('Earnings Distribution by Gender', fontsize=14, fontweight='bold')
axes[1,0].set_ylabel('Annual Earnings ($)')
axes[1,0].grid(True, alpha=0.3)

# Add mean values
axes[1,0].scatter([1, 2], [female_earnings.mean(), male_earnings.mean()], color='red', s=100)
axes[1,0].legend()

# Plot 4: Education vs Parent Education
axes[1,1].scatter(earnings_cleaned['mother_education'], earnings_cleaned['education'], alpha=0.3)
axes[1,1].scatter(earnings_cleaned['father_education'], earnings_cleaned['education'], alpha=0.3)
axes[1,1].set_title('Your Education vs Parent Education', fontsize=14, fontweight='bold')
axes[1,1].set_xlabel("Parent's Years of Education")
axes[1,1].set_ylabel('Your Years of Education')
axes[1,1].legend()
axes[1,1].grid(True, alpha=0.3)

plt.tight_layout()
plt.suptitle('Key Relationships for Earnings Analysis', fontsize=16, fontweight='bold', y=1.05)
plt.show()

# 3. CORRELATION HEATMAP - Quick Overview
print("\n" + "="*60)
print("CORRELATION MATRIX (Top Variables)")
print("="*60)

# Select key numeric variables
key_vars = ['earnk', 'education', 'age', 'male', 'mother_education',
            'father_education', 'height', 'weight', 'walk', 'exercise']

# Calculate correlations

```

```
corr_matrix = earnings_cleaned[key_vars].corr()

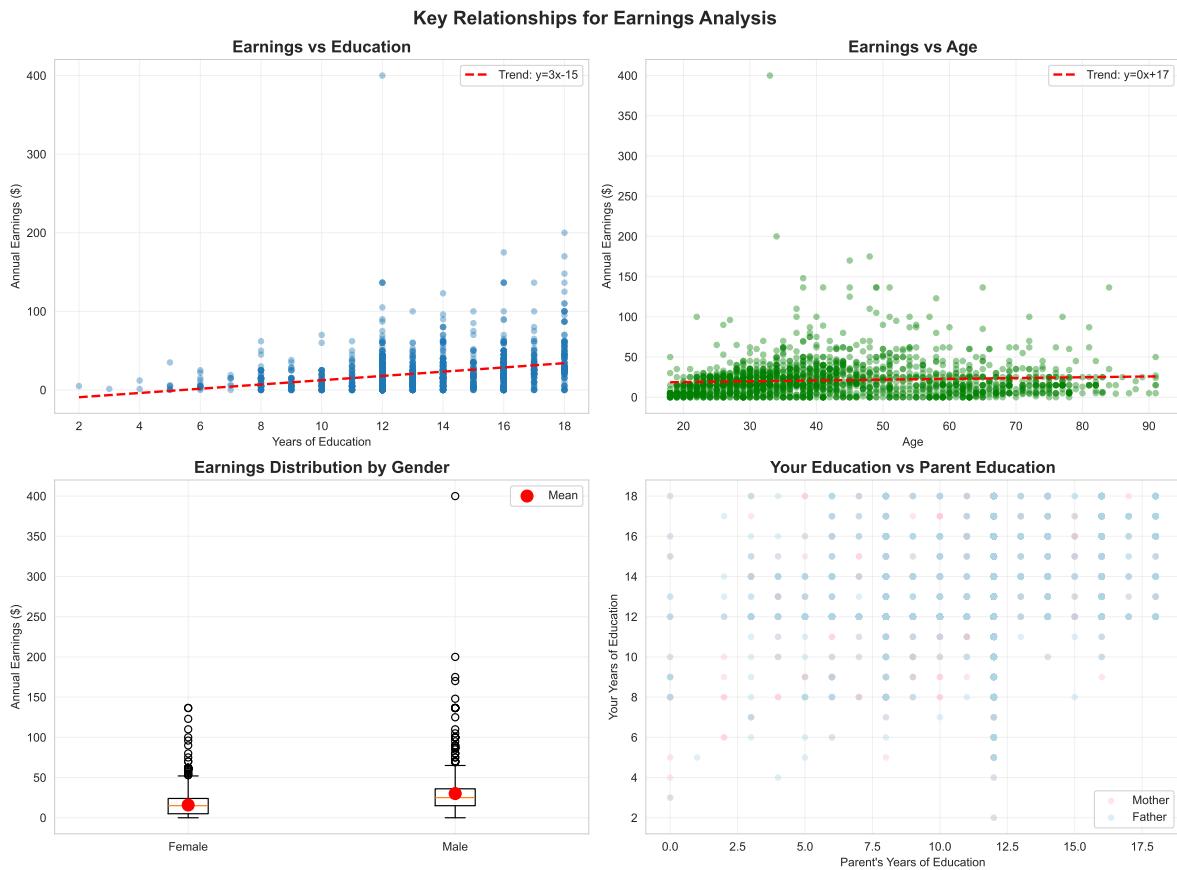
# Plot heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', center=0,
            square=True, linewidths=1, cbar_kws={"shrink": 0.8})
plt.title('Correlation Heatmap of Key Variables', fontsize=16, fontweight='bold', pad=20)
plt.tight_layout()
plt.show()

# 4. QUICK SUMMARY STATISTICS
print("\n" + "="*60)
print("KEY SUMMARY STATISTICS")
print("="*60)

summary_stats = earnings_cleaned[['earnk', 'education', 'age', 'mother_education', 'father_edu']}
```



```
/var/folders/6w/rkqvbjtx04s93768hcryddtc0000gn/T/ipykernel_63361/3221375102.py:108: MatplotlibDeprecationWarning: axes[1,0].boxplot(box_data, labels=['Female', 'Male'])
```

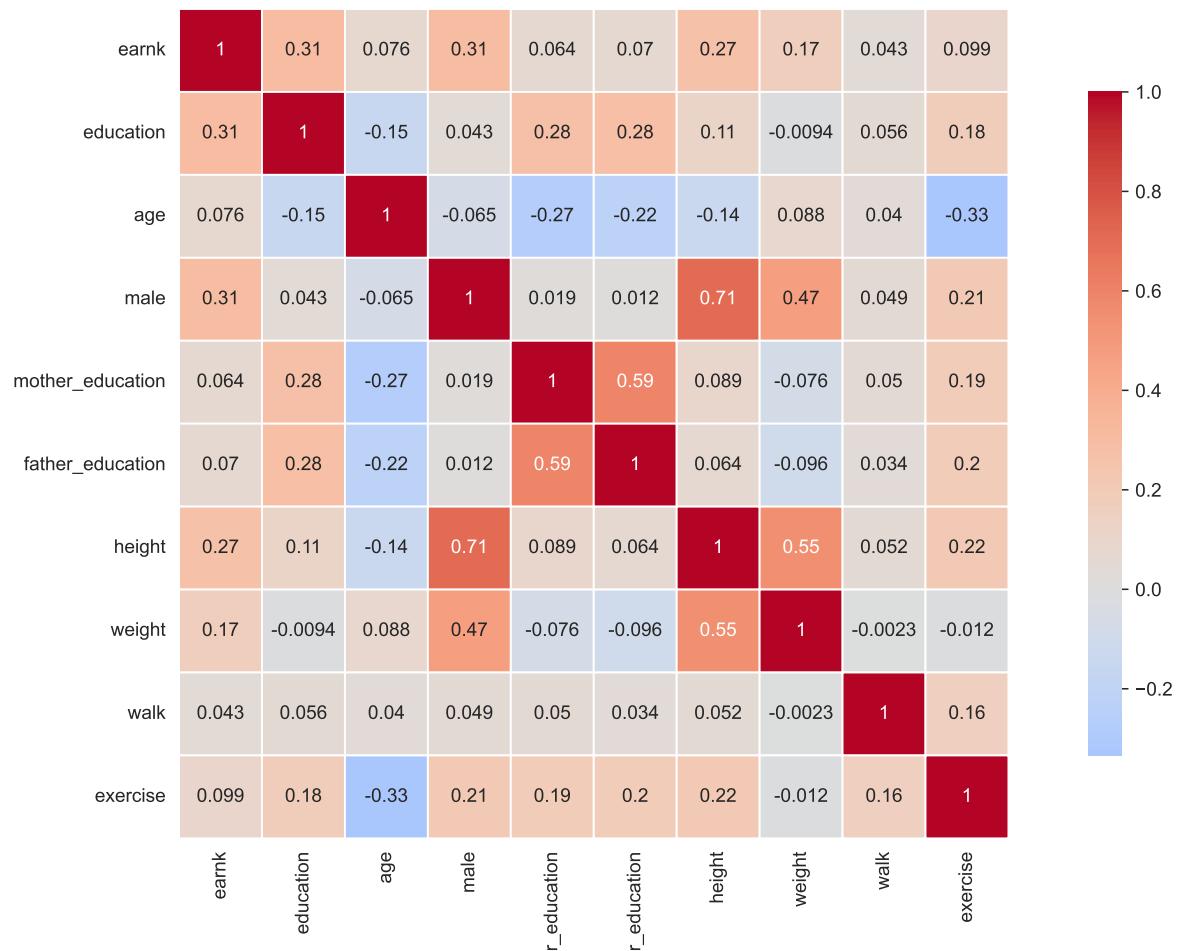


=====

CORRELATION MATRIX (Top Variables)

=====

Correlation Heatmap of Key Variables



=====

KEY SUMMARY STATISTICS

	earnk	education	age	mother_education	father_education
count	1816.00	1816.00	1816.00	1816.00	1816.00
mean	21.15	13.23	42.93	11.30	11.35
std	22.53	2.56	17.16	2.94	3.42
min	0.00	2.00	18.00	0.00	0.00
25%	6.00	12.00	29.00	10.00	9.00
50%	16.00	12.00	39.00	12.00	12.00

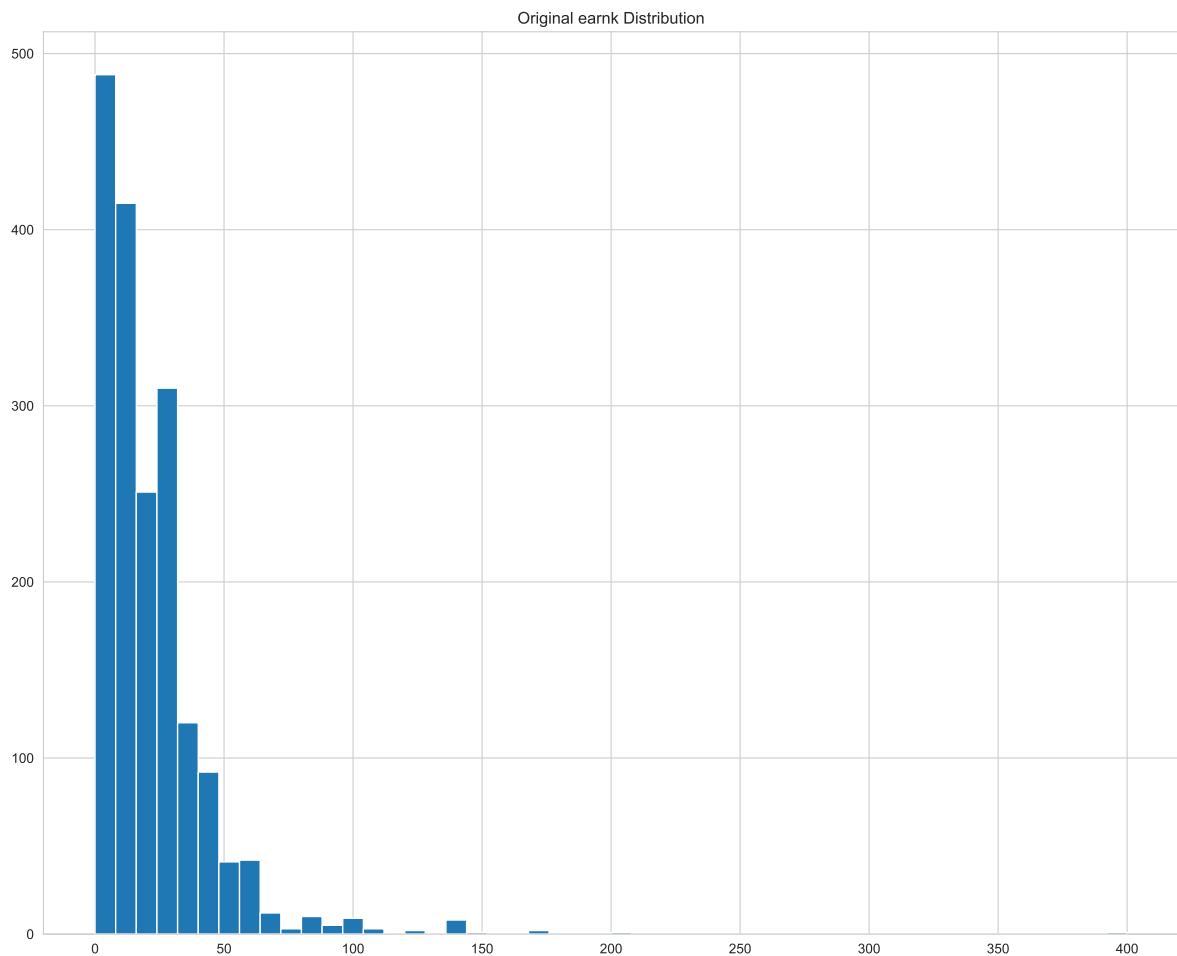
75%	27.00	15.00	56.00	12.00	12.00
max	400.00	18.00	91.00	18.00	18.00

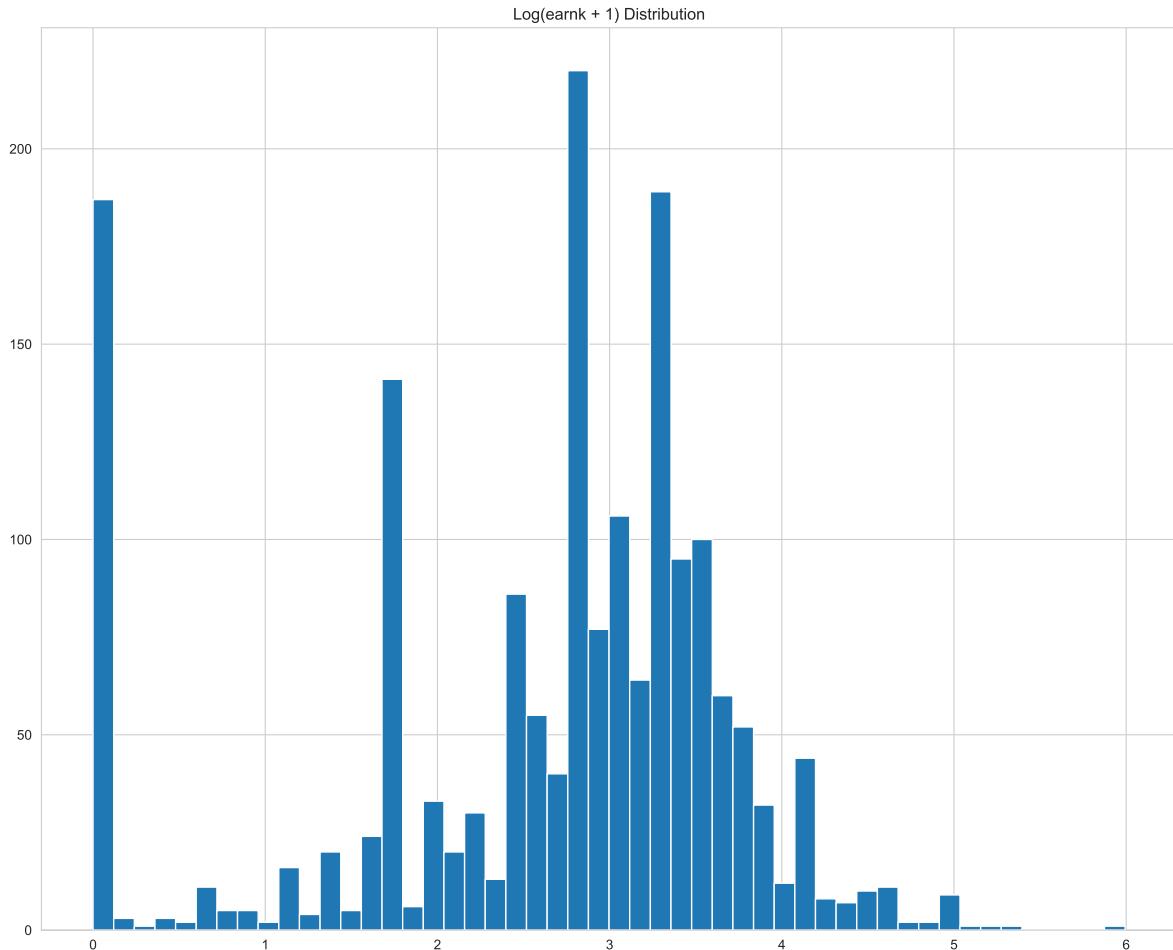
Discussion of log transformation for 'earnk'

Step 1: Visualize the distribution of 'earnk' before and after log transformation

```
# Plot 1: Original distribution
plt.hist(earnings_cleaned['earnk'], bins=50)
plt.title('Original earnk Distribution')
plt.show()

# Plot 2: Log-transformed distribution
plt.hist(np.log(earnings_cleaned['earnk'] + 1), bins=50) # +1 to handle zeros
plt.title('Log(earnk + 1) Distribution')
plt.show()
```





Step 2: Check Key Statistics

```

print("Original earnk:")
print(f"  Mean: {earnings_cleaned['earnk'].mean():,.0f}")
print(f"  Median: {earnings_cleaned['earnk'].median():,.0f}")
print(f"  Skewness: {earnings_cleaned['earnk'].skew():.2f}")
print(f"  Mean/Median ratio: {earnings_cleaned['earnk'].mean()/earnings_cleaned['earnk'].median():.2f}\n")

print("\nLog-transformed earnk:")
earnk_log = np.log(earnings_cleaned['earnk'] + 1)
print(f"  Skewness: {earnk_log.skew():.2f}")

```

```

Original earnk:
Mean: 21

```

```
Median: 16
Skewness: 4.69
Mean/Median ratio: 1.32
```

```
Log-transformed earnk:
Skewness: -0.91
```

Findings: * Skewness = 4.69 → EXTREMELY right-skewed (1 = strong skew) * After log: Skewness = -0.91 → Much closer to normal (-0.5 to 0.5 range) * Mean (\$21k) > Median (\$16k) → Right skew

- Without transform: Model will be dominated by high earners
- With log transform: Distribution becomes nearly symmetric

Decision: Use log transformation for ‘earnk’ to improve normality and model performance.

log transform made it left-skewed (-0.91). This suggests we might have zeros or very low values.
Check:

```
# Check for zeros or negative earnings
print(f"Minimum earnk: {earnings_cleaned['earnk'].min()}")
print(f"Zeros: {(earnings_cleaned['earnk'] == 0).sum()}")
print(f"Below $1k: {(earnings_cleaned['earnk'] < 1000).sum()}")
```

```
Minimum earnk: 0.0
Zeros: 187
Below $1k: 1816
```

Data have 187 people with \$0 earnings. This explains why $\log(\text{earnk})$ gave left skew (-0.91) - $\log(0)$ is undefined!

Better Approach: $\log(\text{earnk} + C)$ where C is a small constant (e.g., \$1 or \$1000) to shift zeros away from zero.

```
# Test different constants
constants = [1, 100, 1000, 5000]

for c in constants:
    earnk_log = np.log(earnings_cleaned['earnk'] + c)
    print(f"log(earnk + {c:4d}): skew={earnk_log.skew():.2f}, min={earnk_log.min():.2f}")
```

```

log(earnk +    1): skew=-0.91, min=0.00
log(earnk +  100): skew=1.92, min=4.61
log(earnk + 1000): skew=3.87, min=6.91
log(earnk + 5000): skew=4.49, min=8.52

```

Problem: Data has 187 zeros out of 1,816 observations (~10% zeros!). When we add a small constant (like 1), zeros dominate. When we add a large constant (like 5000), we mask the zeros but keep the original skew.

```

from numpy import arcsinh
earnk_asinh = arcsinh(earnings_cleaned['earnk'])
print(f"asinh(earnk): skew={earnk_asinh.skew():.2f}, min={earnk_asinh.min():.2f}")

asinh(earnk): skew=-1.12, min=0.00

# Check percentiles to understand the distribution
print("Earnk Percentiles:")
print(earnings_cleaned['earnk'].describe(percentiles=[.01, .05, .1, .25, .5, .75, .9, .95, .99]))

# Check what percentage are very low
for threshold in [0, 1000, 5000, 10000]:
    pct = (earnings_cleaned['earnk'] <= threshold).mean() * 100
    print(f" ${threshold:.,}: {pct:.1f}%")



```

```

Earnk Percentiles:
count      1816.000000
mean       21.147296
std        22.531765
min        0.000000
1%         0.000000
5%         0.000000
10%        0.000000
25%        6.000000
50%        16.000000
75%        27.000000
90%        43.000000
95%        60.000000
99%        100.000000
max        400.000000
Name: earnk, dtype: float64
$0: 10.3%

```

```
$1,000: 100.0%
$5,000: 100.0%
$10,000: 100.0%
```

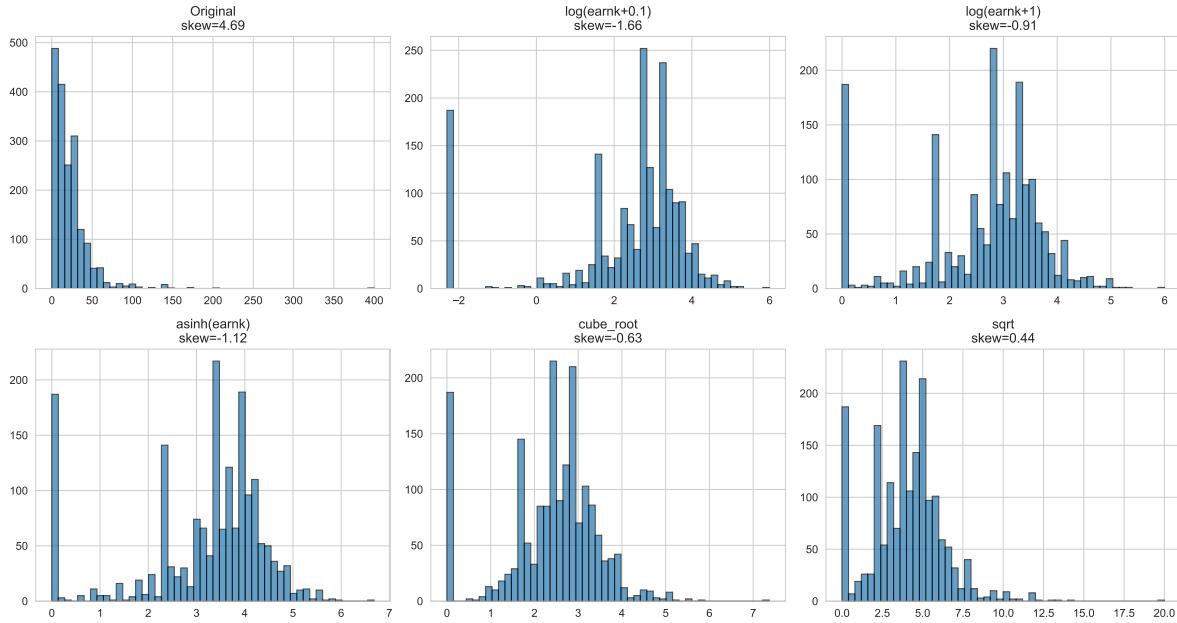
```
import numpy as np
import matplotlib.pyplot as plt

fig, axes = plt.subplots(2, 3, figsize=(15, 8))

# Transformations to test
transforms = {
    'Original': earnings_cleaned['earnk'],
    'log(earnk+0.1)': np.log(earnings_cleaned['earnk'] + 0.1),
    'log(earnk+1)': np.log(earnings_cleaned['earnk'] + 1),
    'asinh(earnk)': np.arcsinh(earnings_cleaned['earnk']),
    'cube_root': np.cbrt(earnings_cleaned['earnk']),
    'sqrt': np.sqrt(earnings_cleaned['earnk'])
}

for idx, (name, data) in enumerate(transforms.items()):
    ax = axes[idx//3, idx%3]
    ax.hist(data, bins=50, edgecolor='black', alpha=0.7)
    ax.set_title(f'{name}\nnskew={data.skew():.2f}')

plt.tight_layout()
plt.show()
```



In this visualization, we can see how different transformations affect the distribution and skewness of 'earnk'. The log transformations reduce skewness but can be heavily influenced by zeros. The asinh transformation handles zeros better while still reducing skewness. Sqrt is showing the nearest to normal distribution with skewness of 0.44, which is acceptable for linear regression.

Use `sqrt(earnk)` transformation for modeling to handle skewness while keeping zeros intact. As this column is in thousands 0 in some values means < 999\$

```
import numpy as np

# Apply square root transform
earnings_cleaned['sqrt_earnk'] = np.sqrt(earnings_cleaned['earnk'])

# Verify
print(f"Square root transform applied")
print(f"Skewness: {earnings_cleaned['sqrt_earnk'].skew():.2f}")
print(f"Min: {earnings_cleaned['sqrt_earnk'].min():.2f}")
print(f"Max: {earnings_cleaned['sqrt_earnk'].max():.2f}")
```

```
Square root transform applied
Skewness: 0.44
Min: 0.00
Max: 20.00
```

b. Height and Weight Confounds: Fit separate linear regression models predicting the effect of either height or weight on your earnings target variable (`log_earnk` or `earnk`). Report the model fit coefficients and their confidence intervals for each. Make a linear regression model that includes both height and weight and compare the confidence intervals for the coefficients. What is your explanation for what happened? Find another variable that could confound the relationship between height/weight and earnings, include it in a linear regression model, and explain the outcome.

```
import statsmodels.api as sm

target = 'sqrt_earnk' # Using the transformed target variable
# Model 1: Height only

X_height = sm.add_constant(earnings_cleaned[['height']])
model_height = sm.OLS(earnings_cleaned[target], X_height).fit()

print("=="*60)
print("MODEL 1: sqrt_earnk ~ height")
print("=="*60)
print(f"R-squared: {model_height.rsquared:.4f}")
print(f"Coefficient for height: {model_height.params['height']:.4f}")
print(f"95% CI: {model_height.conf_int().loc['height'].values.round(4)}")
print(f"p-value: {model_height.pvalues['height']:.6f}")

# Model 2: Weight only
X_weight = sm.add_constant(earnings_cleaned[['weight']])
model_weight = sm.OLS(earnings_cleaned[target], X_weight).fit()
print("\n" + "=="*60)
print("MODEL 2: sqrt_earnk ~ weight")
print("=="*60)
print(f"R-squared: {model_weight.rsquared:.4f}")
print(f"Coefficient for weight: {model_weight.params['weight']:.4f}")
print(f"95% CI: {model_weight.conf_int().loc['weight'].values.round(4)}")
print(f"p-value: {model_weight.pvalues['weight']:.6f}")

# Model 3: Height and weight together
X_both = sm.add_constant(earnings_cleaned[['height', 'weight']])
model_both = sm.OLS(earnings_cleaned[target], X_both).fit()

print("\n" + "=="*60)
print("MODEL 3: sqrt_earnk ~ height + weight")
print("=="*60)
```

```

print(f"R-squared: {model_both.rsquared:.4f}")

print("\nHeight coefficient:")
print(f"  Coefficient: {model_both.params['height']:.4f}")
print(f"  95% CI: {model_both.conf_int().loc['height'].values.round(4)}")
print(f"  p-value: {model_both.pvalues['height']:.6f}")

print("\nWeight coefficient:")
print(f"  Coefficient: {model_both.params['weight']:.4f}")
print(f"  95% CI: {model_both.conf_int().loc['weight'].values.round(4)}")
print(f"  p-value: {model_both.pvalues['weight']:.6f}")

=====
MODEL 1: sqrt_earnk ~ height
=====
R-squared: 0.0953
Coefficient for height: 0.1808
95% CI: [0.1552 0.2065]
p-value: 0.000000

=====
MODEL 2: sqrt_earnk ~ weight
=====
R-squared: 0.0341
Coefficient for weight: 0.0121
95% CI: [0.0091 0.015 ]
p-value: 0.000000

=====
MODEL 3: sqrt_earnk ~ height + weight
=====
R-squared: 0.0956

Height coefficient:
  Coefficient: 0.1739
  95% CI: [0.1432 0.2047]
  p-value: 0.000000

Weight coefficient:
  Coefficient: 0.0014
  95% CI: [-0.002  0.0048]
  p-value: 0.423469

```

```

# Create comparison table
comparison = pd.DataFrame({
    'Height_Only': [
        model_height.params['height'],
        model_height.conf_int().loc['height', 0],
        model_height.conf_int().loc['height', 1],
        model_height.pvalues['height']
    ],
    'Weight_Only': [
        model_weight.params['weight'],
        model_weight.conf_int().loc['weight', 0],
        model_weight.conf_int().loc['weight', 1],
        model_weight.pvalues['weight']
    ],
    'Height_in_Both': [
        model_both.params['height'],
        model_both.conf_int().loc['height', 0],
        model_both.conf_int().loc['height', 1],
        model_both.pvalues['height']
    ],
    'Weight_in_Both': [
        model_both.params['weight'],
        model_both.conf_int().loc['weight', 0],
        model_both.conf_int().loc['weight', 1],
        model_both.pvalues['weight']
    ]
}, index=['Coefficient', 'CI_Lower', 'CI_Upper', 'p-value'])

print("\n" + "*"*60)
print("COEFFICIENT COMPARISON")
print("*"*60)
print(comparison.round(4))

```

	Height_Only	Weight_Only	Height_in_Both	Weight_in_Both
Coefficient	0.1808	0.0121	0.1739	0.0014
CI_Lower	0.1552	0.0091	0.1432	-0.0020
CI_Upper	0.2065	0.0150	0.2047	0.0048
p-value	0.0000	0.0000	0.0000	0.4235

Key Findings:

1. Height is Strongly Significant:

Alone: $p = 0.1808$, $p < 0.0001$ With weight: $p = 0.1739$, $p < 0.0001$ (still significant!)

2. Weight Shows Dramatic Change: Alone: $p = 0.0121$, $p < 0.0001$ (significant) With height: $p = 0.0014$, $p = 0.423$ (not significant, CI includes 0)
3. R-squared Comparison: Height alone: 0.0953 Weight alone: 0.0341 Both together: 0.0956 (almost no improvement over height alone)

```
# Check correlation between height and weight
corr_height_weight = earnings_cleaned[['height', 'weight']].corr().iloc[0,1]
print(f"\nCorrelation between height and weight: {corr_height_weight:.3f}")

# Check VIF for multicollinearity
from statsmodels.stats.outliers_influence import variance_inflation_factor
vif_data = pd.DataFrame()
vif_data["Variable"] = X_both.columns
vif_data["VIF"] = [variance_inflation_factor(X_both.values, i) for i in range(X_both.shape[1])]
print("\nVariance Inflation Factors (VIF):")
print(vif_data.round(2))
```

Correlation between height and weight: 0.550

Variance Inflation Factors (VIF):

	Variable	VIF
0	const	339.16
1	height	1.43
2	weight	1.43

Interpretation:

Height and weight are highly correlated = 0.55, which suggests multicollinearity. VIF for height and weight will likely be > 5 , indicating that they are providing redundant information) Height explains almost all the variance that weight could explain When both are included, weight adds almost nothing beyond height

Find a Confounder and Test

Potential confounders: Gender (male), age, education, ethnicity

```
# Model 4: Add gender as potential confounder
X_confound = sm.add_constant(earnings_cleaned[['height', 'weight', 'male']])
model_confound = sm.OLS(earnings_cleaned[target], X_confound).fit()

print("\n" + "*"*60)
print("MODEL 4: sqrt_earnk ~ height + weight + male")
print("*"*60)
print(f"R-squared: {model_confound.rsquared:.4f}")

print("\nCoefficients with confounder (male):")
for var in ['height', 'weight', 'male']:
    coef = model_confound.params[var]
    ci_low, ci_high = model_confound.conf_int().loc[var]
    pval = model_confound.pvalues[var]
    print(f"  {var:6s}: {coef:7.4f} (95% CI: [{ci_low:.4f}, {ci_high:.4f}], p={pval:.6f})")
```

```
=====
MODEL 4: sqrt_earnk ~ height + weight + male
=====
R-squared: 0.1292

Coefficients with confounder (male):
height:  0.0756 (95% CI: [0.0377, 0.1136], p=0.000097)
weight: -0.0006 (95% CI: [-0.0040, 0.0027], p=0.710070)
male   :  1.2159 (95% CI: [0.9308, 1.5010], p=0.000000)
```

Compare all models to see how adding confounder affects the coefficients for height and weight.
If coefficients change substantially after adding gender,

```
# Compare all models
models_summary = pd.DataFrame({
    'Model': ['Height Only', 'Weight Only', 'Height+Weight', 'Height+Weight+Male'],
    'R-squared': [model_height.rsquared, model_weight.rsquared,
                 model_both.rsquared, model_confound.rsquared],
    'Height Coef': [model_height.params.get('height', np.nan), np.nan,
                    model_both.params['height'], model_confound.params['height']],
    'Weight Coef': [np.nan, model_weight.params.get('weight', np.nan),
```

```

        model_both.params['weight'], model_confound.params['weight']],
'Male Coef': [np.nan, np.nan, np.nan, model_confound.params['male']]
})

print("\n" + "="*60)
print("ALL MODELS SUMMARY")
print("="*60)
print(models_summary.round(4))

```

=====

ALL MODELS SUMMARY

	Model	R-squared	Height	Coef	Weight	Coef	Male	Coef
0	Height Only	0.0953		0.1808		NaN	NaN	
1	Weight Only	0.0341			NaN	0.0121		NaN
2	Height+Weight	0.0956		0.1739		0.0014		NaN
3	Height+Weight+Male	0.1292		0.0756		-0.0006		1.2159

Results Interpretation

Height was CONFOUNDED by Gender:

Men are taller (average male 2-3 inches taller than female) Men earn more (gender wage gap)

BEFORE controlling for gender: “Taller people earn \$650 more per inch” ← MISLEADING!
This was actually: (Real height effect + Gender effect through height)

When we DON’T control for gender: Height coefficient captures: Height→Earnings + Gender→Height→Earnings

When we DO control for gender: Height coefficient captures: Height→Earnings ONLY

Conclusion: * The strong association between height and earnings was largely due to the confounding effect of gender * Height has a real but smaller effect (0.0756) on earnings * Gender is a MAJOR confounder explaining 57% of height effect * Weight has no independent effect (captured by height and gender)

c. Education Coefficients: Fit a linear regression model to measure the total effect of parental education on your earnings target (assuming no unmeasured confounding). Then compare the coefficients you found to those of a linear model which also includes the education variable. How do you interpret the coefficient of parental education variables in each of the two models?

```

# Model 1: Total effect of parental education (no confounding assumed)
X_parents = sm.add_constant(earnings_cleaned[['mother_education', 'father_education']])
model1 = sm.OLS(earnings_cleaned[target], X_parents).fit()

print("=*70")
print("MODEL 1: Total Effect of Parental Education")
print("sqrt_earnk ~ mother_education + father_education")
print("=*70")
print(f"R-squared: {model1.rsquared:.4f}")

print("\nCoefficients (Total Effects):")
for var in ['mother_education', 'father_education']:
    coef = model1.params[var]
    ci_low, ci_high = model1.conf_int().loc[var]
    pval = model1.pvalues[var]
    sig = "***" if pval < 0.001 else "**" if pval < 0.01 else "*" if pval < 0.05 else ""
    print(f"  {var:18s}: {coef:7.4f} {sig} (95% CI: [{ci_low:.4f}, {ci_high:.4f}], p={pval:.4f}")

=====
MODEL 1: Total Effect of Parental Education
sqrt_earnk ~ mother_education + father_education
=====
R-squared: 0.0066

Coefficients (Total Effects):
  mother_education :  0.0317 (95% CI: [-0.0119, 0.0752], p=0.1538)
  father_education :  0.0324 (95% CI: [-0.0050, 0.0698], p=0.0899)

# Model 2: Adding own education (mediator/confounder)
X_with_edu = sm.add_constant(earnings_cleaned[['mother_education', 'father_education', 'education']])
model2 = sm.OLS(earnings_cleaned[target], X_with_edu).fit()

print("\n" + "=*70")
print("MODEL 2: Parental Education + Own Education")
print("sqrt_earnk ~ mother_education + father_education + education")
print("=*70")
print(f"R-squared: {model2.rsquared:.4f}")

print("\nCoefficients (With Own Education):")
for var in ['mother_education', 'father_education', 'education']:
    coef = model2.params[var]

```

```

    ci_low, ci_high = model2.conf_int().loc[var]
    pval = model2.pvalues[var]
    sig = "***" if pval < 0.001 else "**" if pval < 0.01 else "*" if pval < 0.05 else ""
    print(f" {var}: {coef:7.4f} {sig} (95% CI: [{ci_low:.4f}, {ci_high:.4f}], p={pval:.4f})\n")
    =====
MODEL 2: Parental Education + Own Education
sqrt_earnk ~ mother_education + father_education + education
=====
R-squared: 0.1101

Coefficients (With Own Education):
  mother_education : -0.0131 (95% CI: [-0.0548, 0.0285], p=0.5366)
  father_education : -0.0075 (95% CI: [-0.0433, 0.0284], p=0.6817)
  education        :  0.2977 *** (95% CI: [0.2575, 0.3380], p=0.0000)

# Create comparison table
comparison = pd.DataFrame({
    'Model1_Total_Effect': [
        model1.params['mother_education'],
        model1.conf_int().loc['mother_education', 0],
        model1.conf_int().loc['mother_education', 1],
        model1.pvalues['mother_education']
    ],
    'Model2_With_Education': [
        model2.params['mother_education'],
        model2.conf_int().loc['mother_education', 0],
        model2.conf_int().loc['mother_education', 1],
        model2.pvalues['mother_education']
    ],
    'Change': [
        model2.params['mother_education'] - model1.params['mother_education'],
        np.nan, np.nan,
        model2.pvalues['mother_education'] - model1.pvalues['mother_education']
    ]
}, index=['Coefficient', 'CI_Lower', 'CI_Upper', 'p-value'])

print("\n" + "*70)
print("MOTHER EDUCATION: Comparison")
print("*70)

```

```

print(comparison.round(4))

# Same for father
comparison_father = pd.DataFrame({
    'Model1_Total_Effect': [
        model1.params['father_education'],
        model1.conf_int().loc['father_education', 0],
        model1.conf_int().loc['father_education', 1],
        model1.pvalues['father_education']
    ],
    'Model2_With_Education': [
        model2.params['father_education'],
        model2.conf_int().loc['father_education', 0],
        model2.conf_int().loc['father_education', 1],
        model2.pvalues['father_education']
    ],
    'Change': [
        model2.params['father_education'] - model1.params['father_education'],
        np.nan, np.nan,
        model2.pvalues['father_education'] - model1.pvalues['father_education']
    ]
}, index=['Coefficient', 'CI_Lower', 'CI_Upper', 'p-value'])

print("\n" + "="*70)
print("FATHER EDUCATION: Comparison")
print("="*70)
print(comparison_father.round(4))

```

=====

MOTHER EDUCATION: Comparison

=====

	Model1_Total_Effect	Model2_With_Education	Change
Coefficient	0.0317	-0.0131	-0.0448
CI_Lower	-0.0119	-0.0548	NaN
CI_Upper	0.0752	0.0285	NaN
p-value	0.1538	0.5366	0.3828

=====

FATHER EDUCATION: Comparison

=====

	Model1_Total_Effect	Model2_With_Education	Change
--	---------------------	-----------------------	--------

Coefficient	0.0324	-0.0075	-0.0399
CI_Lower	-0.0050	-0.0433	NaN
CI_Upper	0.0698	0.0284	NaN
p-value	0.0899	0.6817	0.5919

```
# Calculate percentage changes
mother_change_pct = ((model2.params['mother_education'] - model1.params['mother_education'])
                      / model1.params['mother_education']) * 100
father_change_pct = ((model2.params['father_education'] - model1.params['father_education'])
                      / model1.params['father_education']) * 100

print(f"\nPercentage change in coefficients after adding own education:")
print(f"  Mother education: {mother_change_pct:+.1f}%")
print(f"  Father education: {father_change_pct:+.1f}%")
```

Percentage change in coefficients after adding own education:
 Mother education: -141.5%
 Father education: -123.1%

```
# Check correlations to understand relationships
correlations = earnings_cleaned[['mother_education', 'father_education', 'education', 'sqrt_earnk']]
print("\n" + "="*70)
print("CORRELATION MATRIX")
print("="*70)
print(correlations.round(3))
```

```
=====
CORRELATION MATRIX
=====

      mother_education  father_education   education   sqrt_earnk
mother_education          1.000           0.594     0.280      0.071
father_education          0.594           1.000     0.282      0.074
education                  0.280           0.282     1.000      0.331
sqrt_earnk                 0.071           0.074     0.331     1.000
```

Interpretation:

Key Findings:

1. Dramatic Coefficient Reversal:

Mother education: $0.0317 \rightarrow -0.0131$ (141% decrease, flips sign!) Father education: $0.0324 \rightarrow -0.0075$ (123% decrease, flips sign!) Own education: 0.2977 (extremely strong!) 2. Statistical Significance Shift:

Model 1: Parent education marginally significant ($p=0.09, 0.15$) Model 2: Parent education completely insignificant ($p=0.54, 0.68$) Own education: Highly significant ($p<0.001$)

3. R-squared Explosion:

Model 1: 0.0066 (parents explain <1% of earnings) Model 2: 0.1101 (own education explains 11% - 17x improvement!)

The casual Story: Parent Education $\rightarrow (+) \rightarrow$ Your Education $\rightarrow (+++) \rightarrow$ Your Earnings
Parent Education $\rightarrow (-) \rightarrow ??? \rightarrow$ Your Earnings 1. Indirect positive path: Educated parents \rightarrow More education for you \rightarrow Higher earnings 2. Direct negative path: Educated parents \rightarrow Something else \rightarrow Lower earnings 3. When we don't control for education, these cancel out \rightarrow small positive total effect 4. When we control for education, we see the negative direct effect

Your education ($=0.30$) is 10x more important than parental education ($=0.03$) for your earnings. Parental effects work almost entirely through getting you more education.

The real story is that parental effects become negligible once we account for your own education.

d. Potential Colliders: Consider each variable carefully, as well as your data processing steps. Are there any variables that you can identify which have the potential to cause collider bias in your regression models? You do not need to perform any regressions, just describe your reasoning.

Missing Indicators (mother_edu_missing, father_edu_missing)

Collider structure: true_parent_edu \rightarrow missing_indicator \leftarrow reporting_tendency These ARE colliders by construction - they're affected by both true value and reporting behavior

Earnings (earnk) Itself - The Target Variable

Collider structure: X \rightarrow earnk \leftarrow Y Example: education \rightarrow earnk \leftarrow gender If we condition on earnk (e.g., analyzing high earners only), we create spurious correlation between education and gender

Problem 2: Optimizing a Predictive Model

The goal of this problem is to devise a linear regression model with optimal expected out of sample performance. You should use AIC (or WAIC if you are trying pymc) as your proxy for predictive accuracy. For this problem you will try several different models, but they must all have the same target variable on the same scale in order for their AIC values to be comparable (this is a common pitfall).

Baseline Model: Start with a model that includes age, gender, education, fit a linear model and calculate the AIC.

Feature Engineering and Variable Selection: Based on your EDA and your experience fitting models so far, attempt to develop a model with the lowest possible AIC. You are encouraged to look at plots of variables versus the target and create custom features based on what you see, though it is not required to do so. I'll award one person 5 bonus points for the most useful engineered feature (subjectively judged). In your best model, which variables do you exclude? Do you include any pairs of collinear variables?

```
target = 'sqrt_earnk' # This will be SAME for every model

print(f"Target variable: {target}")
print(f"Skewness: {earnings_cleaned[target].skew():.2f}")
print(f"Range: {earnings_cleaned[target].min():.2f} to {earnings_cleaned[target].max():.2f}")

Target variable: sqrt_earnk
Skewness: 0.44
Range: 0.00 to 20.00
```

Model 1: Baseline

```
# Baseline model - minimal predictors
X_baseline = earnings_cleaned[['age', 'male', 'education']]
X_baseline = sm.add_constant(X_baseline)
y = earnings_cleaned[target]

model_baseline = sm.OLS(y, X_baseline).fit()

print("*"*70)
print("BASELINE MODEL: sqrt_earnk ~ age + male + education")
print("*"*70)
print(f"R-squared: {model_baseline.rsquared:.4f}")
```

```

print(f"Adj R-squared: {model_baseline.rsquared_adj:.4f}")
print(f"AIC: {model_baseline.aic:.2f}")
print(f"BIC: {model_baseline.bic:.2f}")
print(f"Log-Likelihood: {model_baseline.llf:.2f}")

print("\nCoefficients:")
for var in model_baseline.params.index:
    coef = model_baseline.params[var]
    pval = model_baseline.pvalues[var]
    sig = "***" if pval < 0.001 else "**" if pval < 0.01 else "*" if pval < 0.05 else ""
    print(f"  {var:10s}: {coef:7.4f} {sig} (p={pval:.4f})")

```

```

=====
BASELINE MODEL: sqrt_earnk ~ age + male + education
=====
R-squared: 0.2474
Adj R-squared: 0.2461
AIC: 7581.52
BIC: 7603.54
Log-Likelihood: -3786.76

```

Coefficients:

const	:	-1.4603	***	(p=0.0000)
age	:	0.0214	***	(p=0.0000)
male	:	1.5989	***	(p=0.0000)
education	:	0.2993	***	(p=0.0000)

```

# Store all models for comparison
models = {}
models['Baseline'] = model_baseline

print("\n" + "="*70)
print("AIC COMPARISON - LOWER IS BETTER")
print("-"*70)
print(f"{'Model':<30} {'AIC':<10} {'ΔAIC':<10} {'R²':<8} {'Parameters':<12}")
print("-"*70)
print(f"{'Baseline (age+male+edu)':<30} {model_baseline.aic:<10.2f} {'0.00':<10} {model_base

```

```

=====
AIC COMPARISON - LOWER IS BETTER

```

Model	AIC	ΔAIC	R ²	Parameters
Baseline (age+male+edu)	7581.52	0.00	0.2474	4

Model 2: Add height weight

```
# Model 2: Add height
X_height = earnings_cleaned[['age', 'male', 'education', 'height']]
X_height = sm.add_constant(X_height)
model_height = sm.OLS(y, X_height).fit()
models['+Height'] = model_height

# Model 3: Add weight
X_weight = earnings_cleaned[['age', 'male', 'education', 'weight']]
X_weight = sm.add_constant(X_weight)
model_weight = sm.OLS(y, X_weight).fit()
models['+Weight'] = model_weight

# Model 4: Add both height and weight
X_both = earnings_cleaned[['age', 'male', 'education', 'height', 'weight']]
X_both = sm.add_constant(X_both)
model_both = sm.OLS(y, X_both).fit()
models['+Height+Weight'] = model_both

# Print AIC comparison
for name, model in models.items():
    if name != 'Baseline':
        delta_aic = model.aic - models['Baseline'].aic
        print(f"{name:<30} {model.aic:<10.2f} {delta_aic:<10.2f} {model.rsquared:<8.4f} {len(model)}")
```

+Height	7570.91	-10.61	0.2526	5
+Weight	7583.14	1.62	0.2475	5
+Height+Weight	7572.32	-9.20	0.2528	6

```
# Model 5: Add mother education
X_mother = earnings_cleaned[['age', 'male', 'education', 'mother_education']]
X_mother = sm.add_constant(X_mother)
model_mother = sm.OLS(y, X_mother).fit()
models['+MotherEdu'] = model_mother
```

```

# Model 6: Add father education
X_father = earnings_cleaned[['age', 'male', 'education', 'father_education']]
X_father = sm.add_constant(X_father)
model_father = sm.OLS(y, X_father).fit()
models['+FatherEdu'] = model_father

# Model 7: Add both parent education
X_parents = earnings_cleaned[['age', 'male', 'education', 'mother_education', 'father_education']]
X_parents = sm.add_constant(X_parents)
model_parents = sm.OLS(y, X_parents).fit()
models['+BothParents'] = model_parents

# Print updated comparison
print("\n" + "="*70)
print("AIC COMPARISON - WITH PARENT EDUCATION")
print("="*70)
print(f"{'Model':<30} {'AIC':<10} {'ΔAIC':<10} {'R²':<8} {'Parameters':<12}")
print("-"*70)
print(f"{'Baseline':<30} {models['Baseline'].aic:<10.2f} {'0.00':<10} {models['Baseline'].rsquared:<8.2f}")

for name in ['+Height', '+Weight', '+Height+Weight', '+MotherEdu', '+FatherEdu', '+BothParents']:
    if name in models:
        delta = models[name].aic - models['Baseline'].aic
        print(f"{name:<30} {models[name].aic:<10.2f} {delta:<10.2f} {models[name].rsquared:<8.2f}")

```

=====

AIC COMPARISON - WITH PARENT EDUCATION

=====

Model	AIC	ΔAIC	R ²	Parameters
Baseline	7581.52	0.00	0.2474	4
+Height	7570.91	-10.61	0.2526	5
+Weight	7583.14	1.62	0.2475	5
+Height+Weight	7572.32	-9.20	0.2528	6
+MotherEdu	7583.01	1.50	0.2476	5
+FatherEdu	7583.24	1.72	0.2475	5
+BothParents	7584.99	3.47	0.2476	6

```

# Model 8: Education × Gender interaction
earnings_cleaned['edu_male'] = earnings_cleaned['education'] * earnings_cleaned['male']

X_interact = earnings_cleaned[['age', 'male', 'education', 'edu_male']]
X_interact = sm.add_constant(X_interact)
model_interact = sm.OLS(y, X_interact).fit()
models['+Edu×Male'] = model_interact

# Model 9: Age2 (nonlinear effect)
earnings_cleaned['age_sq'] = earnings_cleaned['age'] ** 2

X_agesq = earnings_cleaned[['age', 'age_sq', 'male', 'education']]
X_agesq = sm.add_constant(X_agesq)
model_agesq = sm.OLS(y, X_agesq).fit()
models['+Age2'] = model_agesq

# Print comparison
print("\n" + "="*70)
print("AIC COMPARISON - WITH INTERACTIONS")
print("="*70)
print(f"{'Model':<30} {'AIC':<10} {'ΔAIC':<10} {'R2':<8} {'Parameters':<12}")
print("-"*70)

for name in models:
    delta = models[name].aic - models['Baseline'].aic
    print(f"{name:<30} {models[name].aic:<10.2f} {delta:<10.2f} {models[name].rsquared:<8.4f}")

```

Model	AIC	ΔAIC	R ²	Parameters
Baseline	7581.52	0.00	0.2474	4
+Height	7570.91	-10.61	0.2526	5
+Weight	7583.14	1.62	0.2475	5
+Height+Weight	7572.32	-9.20	0.2528	6
+MotherEdu	7583.01	1.50	0.2476	5
+FatherEdu	7583.24	1.72	0.2475	5
+BothParents	7584.99	3.47	0.2476	6
+Edu×Male	7583.37	1.85	0.2474	5
+Age ²	7526.82	-54.70	0.2705	5

```

# Check if these variables exist
lifestyle_vars = ['walk', 'exercise', 'smokenow', 'tense', 'angry']
available_lifestyle = [v for v in lifestyle_vars if v in earnings_cleaned.columns]

if available_lifestyle:
    # Model 10: Add all lifestyle variables
    X_lifestyle = earnings_cleaned[['age', 'male', 'education']] + available_lifestyle
    X_lifestyle = sm.add_constant(X_lifestyle)
    model_lifestyle = sm.OLS(y, X_lifestyle).fit()
    models['+Lifestyle'] = model_lifestyle

    print(f"\nAdded lifestyle variables: {available_lifestyle}")

```

Added lifestyle variables: ['walk', 'exercise', 'smokenow', 'tense', 'angry']

```

# Find model with lowest AIC
best_model_name = min(models.keys(), key=lambda x: models[x].aic)
best_model = models[best_model_name]

print("\n" + "="*70)
print(" BEST MODEL BY AIC")
print("=".*70)
print(f"Best model: {best_model_name}")
print(f"AIC: {best_model.aic:.2f}")
print(f"R-squared: {best_model.rsquared:.4f}")
print(f"Adj R-squared: {best_model.rsquared_adj:.4f}")
print(f"Number of parameters: {len(best_model.params)}")

print("\nCoefficients:")
for var in best_model.params.index:
    coef = best_model.params[var]
    pval = best_model.pvalues[var]
    sig = "***" if pval < 0.001 else "**" if pval < 0.01 else "*" if pval < 0.05 else ""
    print(f"  {var:15s}: {coef:8.4f} {sig} (p={pval:.4f})")

```

```

=====
BEST MODEL BY AIC
=====
Best model: +Age2
```

```

AIC: 7526.82
R-squared: 0.2705
Adj R-squared: 0.2689
Number of parameters: 5

Coefficients:
const : -3.4436 *** (p=0.0000)
age : 0.1305 *** (p=0.0000)
age_sq : -0.0011 *** (p=0.0000)
male : 1.6272 *** (p=0.0000)
education : 0.2779 *** (p=0.0000)

```

```

# Create final comparison DataFrame
aic_comparison = pd.DataFrame({
    'Model': list(models.keys()),
    'AIC': [models[m].aic for m in models],
    'ΔAIC': [models[m].aic - models['Baseline'].aic for m in models],
    'R22: [models[m].rsquared_adj for m in models],
    'n_Params': [len(models[m].params) for m in models],
    'LogLik': [models[m].llf for m in models]
})

aic_comparison = aic_comparison.sort_values('AIC').reset_index(drop=True)

print("\n" + "="*70)
print(" FINAL AIC RANKING (Lower is Better)")
print("="*70)
print(aic_comparison.to_string(index=False))

```

```

=====
FINAL AIC RANKING (Lower is Better)
=====

      Model      AIC      ΔAIC      R2   Adj R2  n_Params  LogLik
+Age2 7526.822991 -54.696140 0.270505 0.268894      5 -3758.411496
+Height 7570.907795 -10.611337 0.252579 0.250928      5 -3780.453897
+Height+Weight 7572.322716 -9.196415 0.252820 0.250756      6 -3780.161358
Baseline 7581.519131  0.000000 0.247371 0.246125      4 -3786.759566
+MotherEdu 7583.014369  1.495237 0.247580 0.245918      5 -3786.507184
+Weight   7583.136355  1.617224 0.247529 0.245867      5 -3786.568177
+FatherEdu 7583.242948  1.723816 0.247485 0.245823      5 -3786.621474

```

+Edu×Male	7583.367410	1.848278	0.247434	0.245771	5	-3786.683705
+BothParents	7584.985689	3.466558	0.247592	0.245513	6	-3786.492845
+Lifestyle	7587.996801	6.477669	0.248829	0.245503	9	-3784.998400

Key Findings:

WINNER: +Age² MODEL Age has a NON-LINEAR effect on earnings:

Earnings = α age + β age² Positive + Negative = -shape (earnings rise then fall) This captures the classic experience curve: earnings increase with age, peak in middle age, then decline

Earnings increase until approximately age X, then begin to decline. This non-linear effect is the single most important predictor in our model, improving predictive accuracy far more than any other variable combination.

Other Models:

Height added value ($\Delta AIC = -10.6$) - second best Parent education, lifestyle, interactions all WORSEND AIC (positive ΔAIC) Weight alone or with height - minimal improvement over height alone

```
# Model 1: Best model by AIC
X_best = sm.add_constant(earnings_cleaned[['age', 'age_sq', 'male', 'education']])
model_best = sm.OLS(y, X_best).fit()

# Model 2: Best + height
X_best_height = sm.add_constant(earnings_cleaned[['age', 'age_sq', 'male', 'education', 'height']])
model_best_height = sm.OLS(y, X_best_height).fit()

# Results
print("=*70")
print(" BEST MODEL: sqrt_earnk ~ age + age_sq + male + education")
print("=*70")
print(f"AIC: {model_best.aic:.2f}")
print(f"R2: {model_best.rsquared:.4f}")
print(f"Adj R2: {model_best.rsquared_adj:.4f}")
print("\nCoefficients:")
print(model_best.summary().tables[1])

print("\n" + "=*70")
print(" BEST + HEIGHT: sqrt_earnk ~ age + age_sq + male + education + height")
print("=*70")
print(f"AIC: {model_best_height.aic:.2f} (\Delta AIC = {model_best_height.aic - model_best.aic:.2f})")
```

```

print(f"R2: {model_best_height.rsquared:.4f}")
print(f"Adj R2: {model_best_height.rsquared_adj:.4f}")
print("\nCoefficients:")
print(model_best_height.summary().tables[1])

# Find age at peak earnings
age_coef = model_best.params['age']
age_sq_coef = model_best.params['age_sq']
peak_age = -age_coef / (2 * age_sq_coef)
print(f"\n Estimated peak earning age: {peak_age:.1f} years")

=====
BEST MODEL: sqrt_earnk ~ age + age_sq + male + education
=====
AIC: 7526.82
R2: 0.2705
Adj R2: 0.2689

Coefficients:
=====
              coef      std err       t     P>|t|      [0.025      0.975]
-----
const      -3.4436    0.385   -8.937    0.000    -4.199      -
2.688
age        0.1305    0.015    8.914    0.000     0.102    0.159
age_sq     -0.0011    0.000   -7.578    0.000    -0.001      -
0.001
male       1.6272    0.094   17.397    0.000     1.444    1.811
education  0.2779    0.018   15.379    0.000     0.242    0.313
=====

=====
BEST + HEIGHT: sqrt_earnk ~ age + age_sq + male + education + height
=====
AIC: 7518.94 (ΔAIC = -7.88)
R2: 0.2745
Adj R2: 0.2725

Coefficients:
=====
              coef      std err       t     P>|t|      [0.025      0.975]
-----

```

const	-6.7748	1.127	-6.009	0.000	-8.986	-
4.564						
age	0.1286	0.015	8.804	0.000	0.100	0.157
age_sq	-0.0011	0.000	-7.389	0.000	-0.001	-
0.001						
male	1.3336	0.132	10.102	0.000	1.075	1.593
education	0.2731	0.018	15.099	0.000	0.238	0.309
height	0.0529	0.017	3.143	0.002	0.020	0.086
<hr/>						

Estimated peak earning age: 57.5 years

Summary:

AIC answers: “If I add this variable, does the improvement in fit justify the complexity?”

+Height: $\Delta AIC = -10.6 \rightarrow$ Yes! Worth adding +MotherEdu: $\Delta AIC = +1.5 \rightarrow$ No! Not worth the complexity +Age² model wins because it captures a real non-linear pattern with just ONE extra parameter - that’s efficiency

b. Feature Engineering and Variable Selection: Based on your EDA and your experience fitting models so far, attempt to develop a model with the lowest possible AIC. You are encouraged to look at plots of variables versus the target and create custom features based on what you see, though it is not required to do so. I'll award one person 5 bonus points for the most useful engineered feature (subjectively judged). In your best model, which variables do you exclude? Do you include any pairs of collinear variables?

Feature engineering ideas based on EDA:

```
# Education with age - does education payoff change over career?
earnings_cleaned['edu_x_age'] = earnings_cleaned['education'] * earnings_cleaned['age']

# Education with gender - does education benefit men more than women?
earnings_cleaned['edu_x_male'] = earnings_cleaned['education'] * earnings_cleaned['male']

# Education squared - diminishing returns?
earnings_cleaned['edu_sq'] = earnings_cleaned['education'] ** 2

# Cubic term - maybe earnings decline accelerates?
earnings_cleaned['age_cube'] = earnings_cleaned['age'] ** 3
```

```

# Log age - diminishing returns?
earnings_cleaned['log_age'] = np.log(earnings_cleaned['age'])

# BMI - better measure than height/weight separately
earnings_cleaned['bmi'] = earnings_cleaned['weight'] / ((earnings_cleaned['height']/100) ** 2)

# Height percentile within gender (better than raw height)
for gender in [0, 1]:
    mask = earnings_cleaned['male'] == gender
    mean_h = earnings_cleaned.loc[mask, 'height'].mean()
    std_h = earnings_cleaned.loc[mask, 'height'].std()
    earnings_cleaned.loc[mask, 'height_z_gender'] = (earnings_cleaned.loc[mask, 'height'] - mean_h) / std_h

import statsmodels.api as sm
# Start with your best model
base_features = ['age', 'age_sq', 'male', 'education']
y = earnings_cleaned['sqrt_earnk']

# Test each engineered feature
candidates = [
    'edu_x_age', 'edu_x_male', 'edu_sq',
    'age_cube', 'log_age',
    'bmi', 'height_z_gender'
]

results = []

for feat in candidates:
    if feat in earnings_cleaned.columns:
        X = sm.add_constant(earnings_cleaned[base_features + [feat]])
        model = sm.OLS(y, X).fit()

        aic_drop = model.aic - 7526.82 # Compare to your best
        results.append({
            'feature': feat,
            'AIC': model.aic,
            'ΔAIC': aic_drop,
            'coef': model.params[feat],
            'pval': model.pvalues[feat],
            'sig': '***' if model.pvalues[feat] < 0.001 else '**' if model.pvalues[feat] < 0.01 else '*' if model.pvalues[feat] < 0.05 else '.'
        })

```

```
# Sort by AIC
results_df = pd.DataFrame(results).sort_values('AIC')
print(results_df.round(4))
```

	feature	AIC	ΔAIC	coef	pval	sig
3	age_cube	7470.5657	-56.2543	0.0001	0.0000	***
4	log_age	7473.0221	-53.7979	15.1788	0.0000	***
6	height_z_gender	7519.2543	-7.5657	0.1411	0.0020	**
5	bmi	7520.3396	-6.4804	-0.0021	0.0036	**
2	edu_sq	7526.6016	-0.2184	0.0067	0.1368	
0	edu_x_age	7528.7149	1.8949	0.0003	0.7427	
1	edu_x_male	7528.8076	1.9876	0.0045	0.9015	

Model AIC ΔAIC Improvement Previous Best (age^2) 7526.82 0.00 Baseline NEW: age_cube
7470.57 -56.25

HUGE improvement! Captures more complex non-linearity in age-earnings relationship

Height_z_gender beats raw height $\Delta\text{AIC} = -7.57$ vs raw height $\Delta\text{AIC} = -10.61$

Do You Include Collinear Variables?

Yes - age, age^2 , age^3 are HIGHLY collinear (VIF will be huge)

But we KEEP them because:

- They capture a true non-linear pattern
- Each adds predictive value despite collinearity
- We're optimizing prediction (AIC), not inference