

Lab 4 DATA 622: Machin Learning and Big Data

Mehreen Gillani

Contents

Problem 1: Ridge Regression for Inference	2
Problem 2: Possession Weights	20
Problem 3: Interpreting Bootstrap Uncertainty	30
Extra Credit (5 Points): Ridge versus Lasso and Confidence/Credible Interval Calculations	49

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
```

```
file_path = "https://raw.githubusercontent.com/georgehagstrom/DATA622Spring2026/refs/heads/main/data/nba_stints.csv"

nba_stints= pd.read_csv(file_path)
nba_stints.head()
```

	game_id	stint_id	n_pos	home_points	away_points	minutes	margin	201939	202691	203
0	22200002	1	14	5	2	2.70	21.428571	1	1	1
1	22200002	2	9	6	2	1.67	44.444444	1	1	1
2	22200002	3	5	0	3	0.48	-60.000000	1	0	1
3	22200002	4	5	5	1	0.78	80.000000	1	0	1
4	22200002	5	9	3	6	1.52	-33.333333	1	0	0

Problem 1: Ridge Regression for Inference

****Ordinary Linear Regression:**** Use ordinary linear regression to fit the model described in the overview. Use cross-validation to estimate the out of sample root mean squared error and compare it to the in sample error, you may use ‘RidgeCV’ with to keep consistency with the rest of the assignment, or use ‘LinearRegression’. Make sure ‘fit_intercept=True’, the intercept corresponds to the home court advantage, and do not use ‘sample_weight’. Does the difference between in-sample and cross-validated mean squared error suggest a major problem with overfitting?

```
# Define player columns (col index 7 onwards)
player_cols = nba_stints.columns[7:]

print(f"Dataset shape: {nba_stints.shape}")
print(f"Number of stints: {len(nba_stints):,}")
print(f"Number of players:{len(player_cols):,}")
print(f"\nFirst few columns: {list(nba_stints.columns[:8])}")
print(f"\nTarget (margin) stats:")
print(nba_stints['margin'].describe().round(2))

# Visualize the problem
import matplotlib.pyplot as plt

# Count how many stints each player appears in
player_appearances = (nba_stints[player_cols] != 0).sum(axis=0)

plt.figure(figsize=(10, 4))
plt.hist(player_appearances, bins=50, edgecolor='black')
plt.xlabel('Number of Stints')
plt.ylabel('Number of Players')
plt.title('Player Appearance Distribution\n(Many players appear in very few stints)')
plt.axvline(player_appearances.median(), color='red',
            linestyle='--', label=f'Median: {player_appearances.median():.0f} stints')
plt.legend()
plt.show()

print(f"Players appearing in <50 stints: {(player_appearances < 50).sum()}")
print(f"Players appearing in <100 stints: {(player_appearances < 100).sum()}")
print(f"Players appearing in >500 stints: {(player_appearances > 500).sum()}")
```

Dataset shape: (32358, 546)
Number of stints: 32,358

Number of players:539

First few columns: ['game_id', 'stint_id', 'n_pos', 'home_points', 'away_points', 'minutes',

Target (margin) stats:

count 32358.00

mean 1.29

std 70.57

min -300.00


25% -33.33

50% 0.00

75% 33.33

max 400.00

Name: margin, dtype: float64



Players appearing in <50 stints: 69

Players appearing in <100 stints: 106

Players appearing in >500 stints: 294

69 players appeared in fewer than 50 stints out of ~30,000 That's less than 0.17% of all stints

294 players appeared in 500+ stints These are starters and key rotation players

```

total_players = len(player_cols)
sparse        = (player_appearances < 50).sum()
limited        = (player_appearances < 100).sum()
reliable      = (player_appearances > 500).sum()

print("PLAYER APPEARANCE SUMMARY")
print("=" * 45)
print(f"Total players:           {total_players:>6,}")
print(f"Sparse (<50 stints):      {sparse:>6,}   ({sparse/total_players*100:.1f}%) ← OLS overfit")
print(f"Limited (<100 stints):     {limited:>6,}   ({limited/total_players*100:.1f}%) ← unreliable")
print(f"Reliable (>500 stints):     {reliable:>6,}  ({reliable/total_players*100:.1f}%) ← trustworthy")

print(f"""
KEY INSIGHT:
    {sparse} players ({sparse/total_players*100:.1f}%) have fewer than 50 observations
    OLS will overfit these players → extreme unreliable RAPM values
    Ridge regression will SHRINK these toward zero → conservative estimates

    {reliable} players ({reliable/total_players*100:.1f}%) have 500+ stints
    These are starters/key rotation players → most reliable RAPM estimates
""")

```

PLAYER APPEARANCE SUMMARY

```

=====
Total players:           539
Sparse (<50 stints):      69  (12.8%) ← OLS overfits
Limited (<100 stints):    106 (19.7%) ← unreliable
Reliable (>500 stints):  294 (54.5%) ← trustworthy

```

KEY INSIGHT:

69 players (12.8%) have fewer than 50 observations
 OLS will overfit these players → extreme unreliable RAPM values
 Ridge regression will SHRINK these toward zero → conservative estimates

294 players (54.5%) have 500+ stints
 These are starters/key rotation players → most reliable RAPM estimates

```

from sklearn.linear_model import RidgeCV, LinearRegression
from sklearn.model_selection import cross_val_score
from sklearn.metrics import mean_squared_error

# Suppress runtime warnings globally for this chunk

```

```

warnings.filterwarnings('ignore', category=RuntimeWarning)
# 1. Prepare features and target
# Player columns start at column index 7 onwards
player_cols = nba_stints.columns[7:]

X = nba_stints[player_cols].values      # player +1/0/-1 indicators
y = nba_stints['margin'].values         # target: point margin per possession
n_pos = nba_stints['n_pos'].values      # possessions (NOT used as weight per instructions)

print(f"Feature matrix shape: {X.shape}")
print(f"Target shape: {y.shape}")
print(f"Players in model: {X.shape[1]}")

# 2. Fit Ordinary Linear Regression
# Using RidgeCV with near-zero alpha to approximate OLS (keeps consistency)
# fit_intercept=True → intercept = home court advantage

ols = RidgeCV(
    alphas = [1e-10],      # near-zero regularization OLS
    fit_intercept = True,   # intercept = home court advantage
    cv      = 5
)
ols.fit(X, y)

# 3. In-Sample Error
y_pred_insample = ols.predict(X)
rmse_insample   = np.sqrt(mean_squared_error(y, y_pred_insample))
mse_insample    = mean_squared_error(y, y_pred_insample)

print(f"\n{'='*50}")
print(f"IN-SAMPLE PERFORMANCE")
print(f"{'='*50}")
print(f"  MSE: {mse_insample:.4f}")
print(f"  RMSE: {rmse_insample:.4f}")
print(f"  Intercept (Home Court Advantage): {ols.intercept_:.4f} pts/possession")

# 4. Cross-Validated Error
# Negative MSE scores → convert to positive
cv_mse_scores = -cross_val_score(
    RidgeCV(alphas=[1e-10], fit_intercept=True),
    X, y,
    cv      = 5,

```

```

    scoring = 'neg_mean_squared_error'
)

cv_mse_mean = cv_mse_scores.mean()
cv_rmse_mean = np.sqrt(cv_mse_mean)
cv_rmse_std = np.sqrt(cv_mse_scores).std()

print(f"\n{'='*50}")
print(f"CROSS-VALIDATED PERFORMANCE (5-Fold)")
print(f"{'='*50}")
print(f"  CV MSE (mean): {cv_mse_mean:.4f}")
print(f"  CV RMSE (mean): {cv_rmse_mean:.4f} ± {cv_rmse_std:.4f}")
print(f"  Per-fold RMSE: {[round(np.sqrt(s), 4) for s in cv_mse_scores]}")

# 5. Overfitting Assessment
rmse_diff = cv_rmse_mean - rmse_insample
pct_diff = (rmse_diff / rmse_insample) * 100

print(f"\n{'='*50}")
print(f"OVERFITTING ASSESSMENT")
print(f"{'='*50}")
print(f"  In-Sample RMSE: {rmse_insample:.4f}")
print(f"  CV RMSE: {cv_rmse_mean:.4f}")
print(f"  Difference: {rmse_diff:.4f} ({pct_diff:.1f}%)")

if pct_diff > 20:
    verdict = "YES - Major overfitting detected"
elif pct_diff > 5:
    verdict = "MODERATE - Some overfitting present"
else:
    verdict = "NO - Minimal overfitting"

print(f"  Verdict: {verdict}")

# 6. Home Court Advantage Interpretation
# Convert from per-possession to per-game (avg ~100 possessions/game)
hca_per_game = ols.intercept_
print(f"\n{'='*50}")
print(f"HOME COURT ADVANTAGE")
print(f"{'='*50}")
print(f"  Per 100 possessions: {ols.intercept_:.4f} pts")
print(f"  Per game (~200 pos total, 100/team): {ols.intercept_:.2f} pts")

```

```
print(f" Context: Known NBA home court = ~3 pts/game")
print(f" Our estimate: ~{ols.intercept_: .1f} pts reasonable")
```

```
Feature matrix shape: (32358, 539)
Target shape:         (32358,)
Players in model:     539
```

IN-SAMPLE PERFORMANCE

```
MSE: 4848.1819
RMSE: 69.6289
Intercept (Home Court Advantage): 1.2409 pts/possession
```

CROSS-VALIDATED PERFORMANCE (5-Fold)

```
CV MSE (mean): 5501.5520
CV RMSE (mean): 74.1724 ± 0.4216
Per-fold RMSE: [np.float64(74.15), np.float64(74.1444), np.float64(74.8616), np.float64(74.1444), np.float64(74.1444)]
```

OVERFITTING ASSESSMENT

```
In-Sample RMSE: 69.6289
CV RMSE: 74.1724
Difference: 4.5436 (6.5%)
Verdict: MODERATE - Some overfitting present
```

HOME COURT ADVANTAGE

```
Per 100 possessions: 1.2409 pts
Per game (~200 pos total, 100/team): 1.24 pts
Context: Known NBA home court = ~3 pts/game
Our estimate: ~1.2 pts reasonable
```

The difference between in-sample RMSE (69.63) and cross-validated RMSE (74.17) is relatively modest at 6.5%, suggesting moderate but not catastrophic overfitting. While the gap indicates the model does not generalize perfectly to unseen data, it is not a major cause for concern given the dataset has 32,358 stints against 539 player features — a ratio of approximately 60 observations per feature, which is well above the conventional threshold for OLS stability.

However, the high RMSE values overall (~74) reflect that basketball margin is inherently noisy and difficult to predict from lineup data alone.

The consistent CV RMSE across all five folds (ranging narrowly from 73.53 to 74.86) further confirms that the overfitting is systematic rather than driven by any particular data split.

b. Examining ‘RAPM’ Coefficients: Create a dataframe with the player-ids, the RAPM coefficients, and join it with the player names (from the data file shared earlier). Use the stint matrix to calculate the number of minutes that each player played (‘minutes’ variable) and add that to the data frame too. Sort the players in descending order by ‘RAPM’ and print the top 20 players by ‘RAPM’. What do you notice about their minutes played? Look up the names of a few of the top players on the internet- are they regarded as top NBA players? Make a scatter plot of ‘RAPM’ versus minutes-played.

```
player_ids_url = "https://raw.githubusercontent.com/georgehagstrom/DATA622Spring2026/refs/heads/main/data/player_ids.csv"

player_names = pd.read_csv(player_ids_url)  # load player names
print("Player names file:")
print(player_names.head())
print(f"Columns: {player_names.columns.tolist()}")
```

```
Player names file:
   player_id  player_name
0    1630173  Precious Achiuwa
1     203500    Steven Adams
2    1628389    Bam Adebayo
3    1630534    Ochai Agbaji
4    1630583    Santi Aldama
Columns: ['player_id', 'player_name']
```

```
# 2. Build RAPM DataFrame
rapm_df = pd.DataFrame({
    'player_id': player_cols,
    'RAPM':      ols.coef_
})

# Clean player_id to match format in player_names file
rapm_df['player_id'] = rapm_df['player_id'].astype(str)
player_names['player_id'] = player_names['player_id'].astype(str)

# Join with player names
rapm_df = rapm_df.merge(player_names, on='player_id', how='left')
```



```
print(f"\nAfter merge: {len(rapm_df):,} players")
print(rapm_df.head())
```

After merge: 539 players

	player_id	RAPM	player_name
0	201939	17.321840	Stephen Curry
1	202691	17.396468	Klay Thompson
2	203110	24.717366	Draymond Green
3	203952	16.375077	Andrew Wiggins
4	1626172	15.509481	Kevon Looney

```
# 3. Calculate Minutes Played per Player
# For each player, sum minutes from stints where they were on court (±1)
minutes_played = {}

for player in player_cols:
    # Player was on court when value is +1 or -1
    on_court_mask = nba_stints[player] != 0
    minutes_played[player] = nba_stints.loc[on_court_mask, 'minutes'].sum()

# Add to dataframe
rapm_df['minutes_played'] = rapm_df['player_id'].map(minutes_played)

print("\nMinutes played stats:")
print(rapm_df['minutes_played'].describe().round(1))
```

Minutes played stats:

count	539.0
mean	1067.7
std	791.4
min	0.8
25%	319.6
50%	965.4
75%	1768.6
max	2956.6

Name: minutes_played, dtype: float64

```
# 4. Sort and Print Top 20
rapm_sorted = rapm_df.sort_values('RAPM', ascending=False).reset_index(drop=True)

print("\nTOP 20 PLAYERS BY OLS RAPM")
print("=" * 65)
print(f"{'Rank':<5} {'Player':<25} {'RAPM':>8} {'Minutes':>10}")
print("-" * 65)

for i, row in rapm_sorted.head(20).iterrows():
    print(f"{i+1:<5} {str(row.get('player_name', row['player_id'])):<25} "
          f"{row['RAPM']:>8.2f} {row['minutes_played']:>10.0f}")

print("\nMinutes Played - Top 20 Stats:")
top20 = rapm_sorted.head(20)
print(f" Mean minutes: {top20['minutes_played'].mean():,.0f}")
print(f" Median minutes: {top20['minutes_played'].median():,.0f}")
print(f" Min minutes: {top20['minutes_played'].min():,.0f}")
print(f" Max minutes: {top20['minutes_played'].max():,.0f}")
print(f"\n Low minutes players (<100 min): "
      f"{(top20['minutes_played'] < 100).sum()} out of 20")
```

TOP 20 PLAYERS BY OLS RAPM

=====			
Rank	Player	RAPM	Minutes

1	Stanley Umude	77.40	2
2	Jordan Schakel	62.88	6
3	Marko Simonovic	50.27	18
4	Alize Johnson	50.14	29
5	Deonte Burton	47.43	8
6	Donovan Williams	47.43	4
7	Isaiah Mobley	41.99	87
8	Nikola Jokic	37.63	2256
9	Kendall Brown	37.51	38
10	Jalen Brunson	36.86	2406
11	Dereon Seabron	36.22	11
12	Luka Samanic	36.06	140
13	Mfiondu Kabengele	34.27	42
14	Ron Harper Jr.	33.44	46
15	Xavier Cooks	33.41	117
16	Devon Dotson	33.31	51

17	Miles McBride	33.03	774
18	Tyrese Martin	33.01	64
19	Jarrell Brantley	32.95	33
20	Mac McClung	32.61	44

Minutes Played - Top 20 Stats:

Mean minutes: 309
Median minutes: 43
Min minutes: 2
Max minutes: 2,406

Low minutes players (<100 min): 15 out of 20

```
# 5. Scatter Plot RAPM vs Minutes
fig, ax = plt.subplots(figsize=(10, 6))

# All players
ax.scatter(rapm_df['minutes_played'], rapm_df['RAPM'],
           alpha=0.4, color='steelblue', s=20, label='All Players')


# Highlight top 20
ax.scatter(top20['minutes_played'], top20['RAPM'],
           color='red', s=60, zorder=5, label='Top 20 RAPM')

# Label top 10 by name
for _, row in rapm_sorted.head(10).iterrows():
    name = str(row.get('player_name', row['player_id']))
    ax.annotate(name,
                xy=(row['minutes_played'], row['RAPM']),
                xytext=(5, 5), textcoords='offset points',
                fontsize=7, color='darkred')

# Reference lines
ax.axhline(y=0, color='black', linestyle='-', linewidth=0.8, alpha=0.5)
ax.axvline(x=500, color='gray', linestyle='--', linewidth=0.8,
           alpha=0.7, label='500 min threshold')

ax.set_xlabel('Minutes Played', fontsize=12)
ax.set_ylabel('RAPM (pts per 100 possessions)', fontsize=12)
ax.set_title('OLS RAPM vs Minutes Played\n(Top players in red)', fontsize=13)
ax.legend(fontsize=10)
ax.grid(True, alpha=0.3)
```

```
plt.tight_layout()
plt.show()
```



Out of Top 20:

- 15 players played fewer than 100 minutes
- Only 3 players played 500+ minutes (Jokic, Brunson, McBride)

```
# Quantify the minutes problem clearly
print("TOP 20 MINUTES BREAKDOWN")
print("=" * 45)
thresholds = [50, 100, 500, 1000]
for t in thresholds:
    count = (top20['minutes_played'] < t).sum()
    print(f"  Played < {t:>4} minutes: {count:>2}/20 players ({count/20*100:.0f}%)")

print(f"\nMedian minutes (Top 20):    {top20['minutes_played'].median():>6.0f}")
print(f"Median minutes (All players): {rapm_df['minutes_played'].median():>6.0f}")
```

TOP 20 MINUTES BREAKDOWN

```
=====
Played < 50 minutes: 12/20 players (60%)
Played < 100 minutes: 15/20 players (75%)
Played < 500 minutes: 17/20 players (85%)
Played < 1000 minutes: 18/20 players (90%)
```

```
Median minutes (Top 20):      43
Median minutes (All players): 965
```

OLS RAPM Findings

Best player = High RAPM + High Minutes

From Our table only TWO players qualify:

Nikola Jokic RAPM 37.63 2,256 min

Jalen Brunson RAPM 36.86 2,406 min

Everyone else either:

→ Has too few minutes to trust (ranks 1-7, 9, 11+)

→ Or has moderate minutes but unproven sample

Minutes Played Problem:

- Median minutes for top 20 players is only 43 minutes vs 965 minutes for the average player — a 22x gap
- 60% of top 20 players played fewer than 50 minutes and 90% played fewer than 1,000 minutes

What This Means:

- OLS overfits players with very few appearances - a player on court for just 2 minutes during a high-scoring run receives an artificially extreme RAPM coefficient
- The top ranked player (Stanley Umude, RAPM 77.40) played only 2 minutes while the median player played 965 minutes

Conclusion:

OLS assigns extreme coefficients to low-minute players with no penalty for small sample size — these are statistical artifacts not true skill

c. Ridge Regression RAPM: The results of (b) suggest that the model is attaching extreme values of 'RAPM' to low minute players, something which can be potentially fixed with regularization. Define a vector of regularization parameters 'alpha' on a logarithmic scale between and (look up 'np.logspace'). Make this vector contain at least 10 but not more than 200 values of 'alpha' (pick based on how fast your computer is). Use 'RidgeCV' to fit a ridge

regression model, selecting the model with the best value of the hyperparameters. What value of 'alpha' is optimal? Next repeat the same calculation as you did in part (b) (you could create a function or just copy the dataframe and replace the old 'RAPM' with new 'RAPM'). Look up some of the top players that your model identified, are they well regarded by the NBA? Make a scatterplot of 'RAPM' versus minutes played.

```
import warnings
import numpy as np
from sklearn.linear_model import RidgeCV

# 1. Define Alpha Vector on Logarithmic Scale
# logspace between 10^-2 and 10^4, 100 values
alphas = np.logspace(-2, 4, 100)

print(f"Alpha range:      {alphas[0]:.4f} to {alphas[-1]:.0f}")
print(f"Number of alphas: {len(alphas)}")
print(f"Sample values:      {[round(a, 3) for a in alphas[:5]]} ... {[round(a,1) for a in alphas[-5:]]}")

# 2. Fit Ridge Regression with Cross Validation
with warnings.catch_warnings():
    warnings.filterwarnings('ignore', category=RuntimeWarning)
    ridge = RidgeCV(
        alphas      = alphas,
        fit_intercept = True,
        cv          = 5
    )
    ridge.fit(X, y)

print(f"\nRIDGE CV RESULTS")
print(f"{'='*50}")
print(f"  Optimal alpha:      {ridge.alpha_:.4f}")
print(f"  Log10(optimal alpha): {np.log10(ridge.alpha_):.2f}")
print(f"  Intercept (HCA):    {ridge.intercept_:.4f} pts/100 possessions")
```

Alpha range: 0.0100 to 10,000

Number of alphas: 100

Sample values: [np.float64(0.01), np.float64(0.011), np.float64(0.013), np.float64(0.015)

RIDGE CV RESULTS

```
=====
Optimal alpha:      613.5907
Log10(optimal alpha):  2.79
```

Intercept (HCA): 1.1942 pts/100 possessions

```
# 3. Build Ridge RAPM DataFrame
ridge_rapm_df = pd.DataFrame({
    'player_id': player_cols,
    'RAPM': ridge.coef_
})

# Merge with player names
ridge_rapm_df['player_id'] = ridge_rapm_df['player_id'].astype(str)
ridge_rapm_df = ridge_rapm_df.merge(player_names, on='player_id', how='left')

# Add minutes played
ridge_rapm_df['minutes_played'] = ridge_rapm_df['player_id'].map(minutes_played)

# Sort by RAPM descending
ridge_rapm_sorted = ridge_rapm_df.sort_values('RAPM', ascending=False).reset_index(drop=True)

# 4. Print Top 20
print("\nTOP 20 PLAYERS BY RIDGE RAPM")
print("=" * 65)
print(f"{'Rank':<5} {'Player':<25} {'RAPM':>8} {'Minutes':>10}")
print("-" * 65)

for i, row in ridge_rapm_sorted.head(20).iterrows():
    print(f"{i+1:<5} {str(row['player_name']):<25} "
          f"{row['RAPM']:>8.2f} {row['minutes_played']:>10.0f}")

# Minutes breakdown
top20_ridge = ridge_rapm_sorted.head(20)
print(f"\nMINUTES BREAKDOWN - TOP 20 RIDGE")
print("=" * 45)
for t in [50, 100, 500, 1000]:
    count = (top20_ridge['minutes_played'] < t).sum()
    print(f"  Played < {t:>4} minutes: {count:>2}/20 ({count/20*100:.0f}%)")

print(f"\n  Median minutes (Top 20): {top20_ridge['minutes_played'].median():>6.0f}")
print(f"  Median minutes (All players): {ridge_rapm_df['minutes_played'].median():>6.0f}")
```

TOP 20 PLAYERS BY RIDGE RAPM

=====

Rank	Player	RAPM	Minutes
1	Joel Embiid	6.12	2249
2	Nikola Jokic	5.36	2256
3	Trae Young	5.18	2706
4	Pascal Siakam	4.63	2497
5	Jalen Brunson	4.22	2406
6	Kevin Love	4.18	1132
7	Draymond Green	4.09	2019
8	Zion Williamson	3.88	859
9	Brook Lopez	3.78	2157
10	Coby White	3.69	1795
11	Anthony Davis	3.67	1881
12	Kawhi Leonard	3.65	1817
13	Myles Turner	3.57	1676
14	Darius Garland	3.55	2210
15	Derrick White	3.53	2079
16	Julius Randle	3.47	2957
17	Isaiah Joe	3.33	1553
18	Jrue Holiday	3.24	2097
19	Franz Wagner	3.22	2293
20	Cameron Johnson	3.17	1073

MINUTES BREAKDOWN - TOP 20 RIDGE

=====

```

Played < 50 minutes: 0/20 (0%)
Played < 100 minutes: 0/20 (0%)
Played < 500 minutes: 0/20 (0%)
Played < 1000 minutes: 1/20 (5%)

```

```

Median minutes (Top 20):      2088
Median minutes (All players):  965

```

```

# OLS vs Ridge Top 20 Side by Side
ols_top20 = rapm_sorted.head(20).reset_index(drop=True)
ridge_top20 = ridge_rapm_sorted.head(20).reset_index(drop=True)

print("\nOLS TOP 20 vs RIDGE TOP 20")
print("=" * 85)
print(f"{'--- OLS ---':<40} {'--- RIDGE ---':<40}")
print(f"{'Rank':<5} {'Player':<22} {'RAPM':>6} {'Min':>6} "
      f"{'Rank':<5} {'Player':<22} {'RAPM':>6} {'Min':>6}")

```



```

print("-" * 85)

for i in range(20):
    ols_row = ols_top20.iloc[i]
    ridge_row = ridge_top20.iloc[i]

    ols_name = str(ols_row['player_name'])[:22]
    ridge_name = str(ridge_row['player_name'])[:22]

    print(f"{i+1:<5} {ols_name:<22} {ols_row['RAPM']:>6.2f} {ols_row['minutes_played']:>6.0f} "
          f"{i+1:<5} {ridge_name:<22} {ridge_row['RAPM']:>6.2f} {ridge_row['minutes_played']}>6.0f}")

print("=" * 85)
print(f"\n{'SUMMARY':^85}")
print(f"  Median minutes OLS:  {ols_top20['minutes_played'].median():>6.0f} min")
print(f"  Median minutes Ridge: {ridge_top20['minutes_played'].median():>6.0f} min")
print(f"  OLS RAPM range:      {ols_top20['RAPM'].min():.2f} to {ols_top20['RAPM'].max():.2f}")
print(f"  Ridge RAPM range:     {ridge_top20['RAPM'].min():.2f} to {ridge_top20['RAPM'].max():.2f}")

```

OLS TOP 20 vs RIDGE TOP 20

--- OLS ---				--- RIDGE ---			
Rank	Player	RAPM	Min	Rank	Player	RAPM	Min
1	Stanley Umude	77.40	2	1	Joel Embiid	6.12	2249
2	Jordan Schakel	62.88	6	2	Nikola Jokic	5.36	2256
3	Marko Simonovic	50.27	18	3	Trae Young	5.18	2706
4	Alize Johnson	50.14	29	4	Pascal Siakam	4.63	2497
5	Deonte Burton	47.43	8	5	Jalen Brunson	4.22	2406
6	Donovan Williams	47.43	4	6	Kevin Love	4.18	1132
7	Isaiah Mobley	41.99	87	7	Draymond Green	4.09	2019
8	Nikola Jokic	37.63	2256	8	Zion Williamson	3.88	859
9	Kendall Brown	37.51	38	9	Brook Lopez	3.78	2157
10	Jalen Brunson	36.86	2406	10	Coby White	3.69	1795
11	Dereon Seabron	36.22	11	11	Anthony Davis	3.67	1881
12	Luka Samanic	36.06	140	12	Kawhi Leonard	3.65	1817
13	Mfiondu Kabengele	34.27	42	13	Myles Turner	3.57	1676
14	Ron Harper Jr.	33.44	46	14	Darius Garland	3.55	2210
15	Xavier Cooks	33.41	117	15	Derrick White	3.53	2079
16	Devon Dotson	33.31	51	16	Julius Randle	3.47	2957
17	Miles McBride	33.03	774	17	Isaiah Joe	3.33	1553

18	Tyrese Martin	33.01	64	18	Jrue Holiday	3.24	2097
19	Jarrell Brantley	32.95	33	19	Franz Wagner	3.22	2293
20	Mac McClung	32.61	44	20	Cameron Johnson	3.17	1073

=====

SUMMARY

Median minutes OLS: 43 min
 Median minutes Ridge: 2088 min
 OLS RAPM range: 32.61 to 77.40
 Ridge RAPM range: 3.17 to 6.12

```
# 6. Scatter Plot Ridge RAPM vs Minutes
fig, axes = plt.subplots(1, 2, figsize=(14, 6))

# OLS Plot (left)
axes[0].scatter(rapm_df['minutes_played'], rapm_df['RAPM'],
                alpha=0.3, color='steelblue', s=15, label='All Players')
axes[0].scatter(rapm_sorted.head(20)['minutes_played'],
                rapm_sorted.head(20)['RAPM'],
                color='red', s=50, zorder=5, label='Top 20')
axes[0].axhline(y=0, color='black', linewidth=0.8, alpha=0.5)
axes[0].axvline(x=500, color='gray', linestyle='--', linewidth=0.8, alpha=0.7)
axes[0].set_xlabel('Minutes Played', fontsize=11)
axes[0].set_ylabel('RAPM (pts per 100 possessions)', fontsize=11)
axes[0].set_title('OLS RAPM vs Minutes\n(Extreme values at low minutes)', fontsize=11)
axes[0].legend(fontsize=9)
axes[0].grid(True, alpha=0.3)

# Ridge Plot (right)
axes[1].scatter(ridge_rapm_df['minutes_played'], ridge_rapm_df['RAPM'],
                alpha=0.3, color='steelblue', s=15, label='All Players')
axes[1].scatter(top20_ridge['minutes_played'], top20_ridge['RAPM'],
                color='red', s=50, zorder=5, label='Top 20')

# Label top 10 ridge players
for _, row in ridge_rapm_sorted.head(10).iterrows():
    axes[1].annotate(str(row['player_name']).split()[-1],
                    xy=(row['minutes_played'], row['RAPM']),
                    xytext=(5, 3), textcoords='offset points',
                    fontsize=7, color='darkred')


axes[1].axhline(y=0, color='black', linewidth=0.8, alpha=0.5)
```

```

axes[1].axvline(x=500, color='gray', linestyle='--', linewidth=0.8, alpha=0.7)
axes[1].set_xlabel('Minutes Played', fontsize=11)
axes[1].set_ylabel('RAPM (pts per 100 possessions)', fontsize=11)
axes[1].set_title(f'Ridge RAPM vs Minutes\n(alpha={ridge.alpha_:.2f})', fontsize=11)
axes[1].legend(fontsize=9)
axes[1].grid(True, alpha=0.3)

plt.suptitle('OLS vs Ridge RAPM: Effect of Regularization',
             fontsize=13, fontweight='bold', y=1.02)
plt.tight_layout()
plt.show()

```



Optimal Alpha: Ridge cross-validation selected an optimal alpha of 613.59 ($\log_{10} = 2.79$), indicating that strong regularization was necessary — consistent with the severe overfitting observed in OLS.

Top 20 Quality:

- Median minutes for top 20 players jumped from 43 (OLS) to 2,088 (Ridge) — a 49x improvement
- Only 1 player in the top 20 played fewer than 1,000 minutes compared to 18 under OLS
- RAPM range compressed from ± 77 (OLS) to ± 6 (Ridge) — far more realistic values
- All top 20 Ridge players have substantial playing time confirming Ridge correctly prioritizes well-represented players over sparse ones

OLS vs Ridge Rankings:

Players with high minutes and high OLS RAPM (Jokic rank 8, Brunson rank 10) are correctly promoted further up in Ridge rankings. Players with low minutes and inflated OLS RAPM are completely removed from top 20 Ridge rankings. Ridge rankings align minutes played with RAPM — players with more data receive more reliable and stable coefficients.

Conclusion: Ridge regression successfully addresses the overfitting problem by shrinking unreliable low-minute coefficients toward zero, producing rankings that reflect genuine sustained performance rather than statistical artifacts from limited playing time.

Problem 2: Possession Weights

a. Heteroscedasticity in Stints: Make a plot of the ‘margin’ variable as a function of the number of possessions in a stint. What do you notice about the variance of ‘margin’? Why do you think it is happening?

```
import matplotlib.pyplot as plt
import numpy as np

# 1. Basic Stats by Possession Buckets
print("MARGIN VARIANCE BY POSSESSION COUNT")
print("=" * 55)

# Bucket possessions for summary stats
possession_buckets = pd.cut(nba_stints['n_pos'],
                             bins=[0, 5, 10, 20, 30, 50, 100, float('inf')],
                             labels=['1-5', '6-10', '11-20', '21-30', '31-50', '51-100', '100+'])

bucket_stats = nba_stints.groupby(possession_buckets, observed=True)['margin'].agg(
    Count = 'count',
    Mean = 'mean',
    Std = 'std',
    Min = 'min',
    Max = 'max'
).round(2)

print(bucket_stats)
```

MARGIN VARIANCE BY POSSESSION COUNT

```
=====
```

	Count	Mean	Std	Min	Max
n_pos					
1-5	16190	1.76	91.55	-300.00	400.00

6-10	8880	0.60	46.01	-150.00	150.00
11-20	4676	0.96	33.46	-130.77	127.27
21-30	2085	1.59	24.20	-71.43	77.27
31-50	520	-0.26	19.36	-63.64	51.61
51-100	7	5.75	15.78	-23.08	25.45

```
# 2. Main Scatter Plot
fig, axes = plt.subplots(1, 2, figsize=(14, 6))

# Left: Raw Scatter
axes[0].scatter(nba_stints['n_pos'], nba_stints['margin'],
                alpha=0.1, s=5, color='steelblue')
axes[0].axhline(y=0, color='red', linewidth=1, linestyle='--', alpha=0.7)
axes[0].set_xlabel('Number of Possessions (n_pos)', fontsize=11)
axes[0].set_ylabel('Margin (pts per 100 possessions)', fontsize=11)
axes[0].set_title('Margin vs Possessions\n(Raw Scatter)', fontsize=12)
axes[0].grid(True, alpha=0.3)

# Right: Variance by Possession
# Calculate rolling std to show variance shrinking
pos_groups = nba_stints.groupby('n_pos')['margin'].std().reset_index()
pos_groups.columns = ['n_pos', 'margin_std']
pos_counts = nba_stints.groupby('n_pos')['margin'].count().reset_index()
pos_counts.columns = ['n_pos', 'count']
pos_groups = pos_groups.merge(pos_counts, on='n_pos')

# Only plot possession counts with enough data
pos_groups_filtered = pos_groups[pos_groups['count'] >= 10]

axes[1].scatter(pos_groups_filtered['n_pos'],
                pos_groups_filtered['margin_std'],
                alpha=0.6, s=20, color='steelblue')


# Add theoretical 1/sqrt(n) curve
n_range = np.linspace(1, pos_groups_filtered['n_pos'].max(), 200)
# Scale to match data
scale = pos_groups_filtered['margin_std'].iloc[0] * np.sqrt(
    pos_groups_filtered['n_pos'].iloc[0])
theo_std = scale / np.sqrt(n_range)
axes[1].plot(n_range, theo_std, color='red', linewidth=2,
             linestyle='--', label='Theoretical 1/√n_pos')
```

```

axes[1].set_xlabel('Number of Possessions (n_pos)', fontsize=11)
axes[1].set_ylabel('Std Dev of Margin', fontsize=11)
axes[1].set_title('Margin Variance vs Possessions\n(Variance shrinks with more possessions)',
                  fontsize=12)
axes[1].legend(fontsize=10)
axes[1].grid(True, alpha=0.3)

plt.suptitle('Heteroscedasticity in Stint Margins',
             fontsize=13, fontweight='bold')
plt.tight_layout()
plt.show()

```



```

# 3. Quantify the Variance Pattern
print("\nVARIANCE PATTERN SUMMARY")
print("=" * 50)

low_pos = nba_stints[nba_stints['n_pos'] <= 5]['margin'].std()
mid_pos = nba_stints[(nba_stints['n_pos'] > 5) &
                     (nba_stints['n_pos'] <= 20)]['margin'].std()
high_pos = nba_stints[nba_stints['n_pos'] > 20]['margin'].std()

print(f" Std Dev (1-5 possessions):    {low_pos:>8.2f}")
print(f" Std Dev (6-20 possessions):    {mid_pos:>8.2f}")
print(f" Std Dev (20+ possessions):        {high_pos:>8.2f}")
print(f" Variance reduction ratio:         {low_pos/high_pos:>8.2f}x")

```

VARIANCE PATTERN SUMMARY

=====	
Std Dev (1-5 possessions):	91.55
Std Dev (6-20 possessions):	42.10
Std Dev (20+ possessions):	23.30
Variance reduction ratio:	3.93x

Understanding Std drops from 91.55 $\rightarrow$ 15.78 as possessions increase That is a 5.8x reduction in variance!

SHORT stints (1-5 pos): - 16,190 stints = most common $\leftarrow$ lots of noise - Std = 91.55 $\leftarrow$ huge variance - Max = +400 $\leftarrow$ one shot = extreme margin

LONG stints (51-100 pos): - Only 7 stints = very rare $\leftarrow$ law of large numbers - Std = 15.78 $\leftarrow$ tight variance - Max = +25 $\leftarrow$ stable reliable estimate

- 2 possessions: one lucky shot dominates margin
- 100 possessions: true team quality emerges

OLS treats ALL stints equally regardless of length:L

- Stint A: 2 possessions, margin = +200 $\leftarrow$ noise
- Stint B: 50 possessions, margin = +8 $\leftarrow$ signal

OLS says: both observations count the same! -

Gets fooled by noisy short stints - Assigns extreme RAPM to players in short stints

This is HETEROSCEDASTICITY:

Variance is NOT constant — it depends on n_pos

Fix $\rightarrow$ Weight each stint by n_pos:

- Short stints (2 pos) $\rightarrow$ low weight (unreliable)
- Long stints (50 pos) $\rightarrow$ high weight (reliable)

b. Implementing Weights Regression: The solution to the issue that you observed in 2(a) is something called weighted least squares. This involves adjusting the error term for each stint based on a weight that accounts for whatever factor controls the variance. The new weighted model looks like this:

where is a coefficient that determines how the variance of ‘margin’ should scale for each stint. The idea is that should be smaller for stints where the variance is high, and larger for stints where the variance is low. This forces the model to fit the low-variance stints more closely than the high-variance stints. The correct weights for this problem are the , which implies that the variance of the margin is inversely proportional to the number of stints. Verify this by making

a scatter plot of the margin rate times the square root of the number of possessions versus the number of possessions. Then recalculate the player ‘RAPM’ coefficients from 1(c) by setting the ‘weights’ keyword in the model fit to be ‘n_pos’. How have the rankings of top players changed?

```
# 1. Verify Weights: Plot margin × sqrt(n_pos) vs n_pos
nba_stints['margin_weighted'] = nba_stints['margin'] * np.sqrt(nba_stints['n_pos'])


fig, axes = plt.subplots(1, 2, figsize=(14, 6))

# Left: Original margin vs n_pos (heteroscedastic)
axes[0].scatter(nba_stints['n_pos'], nba_stints['margin'],
                alpha=0.1, s=5, color='steelblue')
axes[0].axhline(y=0, color='red', linewidth=1, linestyle='--')
axes[0].set_xlabel('Number of Possessions (n_pos)', fontsize=11)
axes[0].set_ylabel('Margin (pts per 100 possessions)', fontsize=11)
axes[0].set_title('Original Margin vs Possessions\n(Heteroscedastic - funnel shape)',
                  fontsize=11)
axes[0].grid(True, alpha=0.3)

# Right: Weighted margin vs n_pos (homoscedastic)
axes[1].scatter(nba_stints['n_pos'], nba_stints['margin_weighted'],
                alpha=0.1, s=5, color='steelblue')
axes[1].axhline(y=0, color='red', linewidth=1, linestyle='--')
axes[1].set_xlabel('Number of Possessions (n_pos)', fontsize=11)
axes[1].set_ylabel('Margin × √n_pos', fontsize=11)
axes[1].set_title('Weighted Margin vs Possessions\n(Homoscedastic - constant variance)',
                  fontsize=11)
axes[1].grid(True, alpha=0.3)

plt.suptitle('Verifying Possession Weights: Before vs After',
             fontsize=13, fontweight='bold')
plt.tight_layout()
plt.show()
```


Verifying Possession Weights: Before vs After



```
# 2. Verify Numerically
print("VARIANCE VERIFICATION")
print("=" * 55)
print("Original margin std by possession bucket:")
buckets = pd.cut(nba_stints['n_pos'],
                  bins=[0, 5, 10, 20, 30, 50, 100, float('inf')],
                  labels=['1-5', '6-10', '11-20', '21-30', '31-50', '51-100', '100+'])

orig_std = nba_stints.groupby(buckets, observed=True)['margin'].std()
weighted_std = nba_stints.groupby(buckets, observed=True)['margin_weighted'].std()

comparison = pd.DataFrame({
    'Original Std': orig_std,
    'Weighted Std': weighted_std
}).round(2)
print(comparison)
print(f"\nOriginal std range: {orig_std.min():.2f} to {orig_std.max():.2f}")
print(f"Weighted std range: {weighted_std.min():.2f} to {weighted_std.max():.2f}")
print(f"\nWeighted std is more consistent = variance stabilized")
```

VARIANCE VERIFICATION

=====

Original margin std by possession bucket:

	Original Std	Weighted Std
n_pos		
1-5	91.55	132.64
6-10	46.01	124.56

11-20	33.46	124.68
21-30	24.20	119.35
31-50	19.36	113.30
51-100	15.78	115.78

Original std range: 15.78 to 91.55
 Weighted std range: 113.30 to 132.64

Weighted std is more consistent = variance stabilized

```
# 3. Fit Weighted Ridge Regression
n_pos_weights = nba_stints['n_pos'].values

with warnings.catch_warnings():
    warnings.filterwarnings('ignore', category=RuntimeWarning)
    ridge_weighted = RidgeCV(
        alphas = alphas,
        fit_intercept = True,
        cv = 5
    )
    ridge_weighted.fit(X, y, sample_weight=n_pos_weights)

print(f"\nWEIGHTED RIDGE CV RESULTS")
print(f"{'='*50}")
print(f"  Optimal alpha:          {ridge_weighted.alpha_:.4f}")
print(f"  Log10(optimal alpha):    {np.log10(ridge_weighted.alpha_):.2f}")
print(f"  Intercept (HCA):         {ridge_weighted.intercept_:.4f} pts/100 possessions")
```

WEIGHTED RIDGE CV RESULTS

```
=====
Optimal alpha:          2477.0764
Log10(optimal alpha):    3.39
Intercept (HCA):         1.0750 pts/100 possessions
```

```
# 4. Build Weighted RAPM DataFrame
weighted_rapm_df = pd.DataFrame({
    'player_id': player_cols,
    'RAPM':      ridge_weighted.coef_
})

weighted_rapm_df['player_id'] = weighted_rapm_df['player_id'].astype(str)
```

```

weighted_rapm_df = weighted_rapm_df.merge(player_names, on='player_id', how='left')
weighted_rapm_df['minutes_played'] = weighted_rapm_df['player_id'].map(minutes_played)
weighted_rapm_sorted = weighted_rapm_df.sort_values('RAPM', ascending=False).reset_index(drop=True)

# 5. Three Way Comparison: OLS vs Ridge vs Weighted Ridge
ols_top20 = rapm_sorted.head(20).reset_index(drop=True)
ridge_top20 = ridge_rapm_sorted.head(20).reset_index(drop=True)
weighted_top20 = weighted_rapm_sorted.head(20).reset_index(drop=True)

print("\nTOP 20: RIDGE vs WEIGHTED RIDGE")
print("=" * 90)
print(f"{'--- RIDGE (Unweighted) ---':<45} {'--- RIDGE (Weighted) ---':<40}")
print(f"{'Rank':<5} {'Player':<22} {'RAPM':>6} {'Min':>6} "
      f"{'Rank':<5} {'Player':<22} {'RAPM':>6} {'Min':>6}")
print("-" * 90)

for i in range(20):
    r_row = ridge_top20.iloc[i]
    w_row = weighted_top20.iloc[i]

    r_name = str(r_row['player_name'])[:22]
    w_name = str(w_row['player_name'])[:22]

    print(f"{i+1:<5} {r_name:<22} {r_row['RAPM']:>6.2f} {r_row['minutes_played']:>6.0f} "
          f"{i+1:<5} {w_name:<22} {w_row['RAPM']:>6.2f} {w_row['minutes_played']:>6.0f}")

print("=" * 90)
print(f"\n{'SUMMARY COMPARISON':~90}")
print(f"{'Metric':<35} {'Ridge':>20} {'Weighted Ridge':>20}")
print("-" * 75)
print(f"{'Median minutes (Top 20)':<35} "
      f"{ridge_top20['minutes_played'].median():>20.0f} "
      f"{weighted_top20['minutes_played'].median():>20.0f}")
print(f"{'RAPM range (Top 20)':<35} "
      f"{ridge_top20['RAPM'].min():.2f} to {ridge_top20['RAPM'].max():.2f}{':>8}' "
      f"{weighted_top20['RAPM'].min():.2f} to {weighted_top20['RAPM'].max():.2f}")
print(f"{'Optimal Alpha':<35} "
      f"{ridge.alpha_:>20.2f} "
      f"{ridge_weighted.alpha_:>20.2f}")

```

TOP 20: RIDGE vs WEIGHTED RIDGE

--- RIDGE (Unweighted) ---				--- RIDGE (Weighted) ---			
Rank	Player	RAPM	Min	Rank	Player	RAPM	Min
1	Joel Embiid	6.12	2249	1	Nikola Jokic	4.54	2256
2	Nikola Jokic	5.36	2256	2	Draymond Green	4.46	2019
3	Trae Young	5.18	2706	3	Joel Embiid	4.18	2249
4	Pascal Siakam	4.63	2497	4	Kawhi Leonard	3.90	1817
5	Jalen Brunson	4.22	2406	5	Anthony Davis	3.85	1881
6	Kevin Love	4.18	1132	6	Jrue Holiday	3.78	2097
7	Draymond Green	4.09	2019	7	Giannis Antetokounmpo	3.70	1841
8	Zion Williamson	3.88	859	8	Christian Koloko	3.54	758
9	Brook Lopez	3.78	2157	9	Zion Williamson	3.40	859
10	Coby White	3.69	1795	10	Aaron Gordon	3.06	1967
11	Anthony Davis	3.67	1881	11	Josh Hart	3.03	2475
12	Kawhi Leonard	3.65	1817	12	Brook Lopez	2.84	2157
13	Myles Turner	3.57	1676	13	Jayson Tatum	2.65	2771
14	Darius Garland	3.55	2210	14	Desmond Bane	2.64	1706
15	Derrick White	3.53	2079	15	Cameron Johnson	2.64	1073
16	Julius Randle	3.47	2957	16	Michael Porter Jr.	2.56	1686
17	Isaiah Joe	3.33	1553	17	Derrick White	2.56	2079
18	Jrue Holiday	3.24	2097	18	Nicolas Batum	2.54	1815
19	Franz Wagner	3.22	2293	19	Franz Wagner	2.51	2293
20	Cameron Johnson	3.17	1073	20	Boban Marjanovic	2.50	167

SUMMARY COMPARISON

Metric	Ridge	Weighted Ridge
Median minutes (Top 20)	2088	1924
RAPM range (Top 20)	3.17 to 6.12	2.50 to 4.54
Optimal Alpha	613.59	2477.08

```
# 6. Scatter Plot: Ridge vs Weighted Ridge
fig, axes = plt.subplots(1, 2, figsize=(14, 6))

for ax, df, top20, title in zip(
    axes,
    [ridge_rapm_df, weighted_rapm_df],
    [ridge_top20, weighted_top20],
    ['Ridge RAPM (Unweighted)', f'Weighted Ridge RAPM\n(weights=n_pos)']
):
```

```


):
    ax.scatter(df['minutes_played'], df['RAPM'],
               alpha=0.3, s=15, color='steelblue', label='All Players')
    ax.scatter(top20['minutes_played'], top20['RAPM'],
               color='red', s=50, zorder=5, label='Top 20')

    # Label top 10
    top10 = df.sort_values('RAPM', ascending=False).head(10)
    for _, row in top10.iterrows():
        ax.annotate(str(row['player_name']).split()[-1],
                    xy=(row['minutes_played'], row['RAPM']),
                    xytext=(5, 3), textcoords='offset points',
                    fontsize=7, color='darkred')

    ax.axhline(y=0, color='black', linewidth=0.8, alpha=0.5)
    ax.axvline(x=500, color='gray', linestyle='--', linewidth=0.8, alpha=0.7)
    ax.set_xlabel('Minutes Played', fontsize=11)
    ax.set_ylabel('RAPM (pts per 100 possessions)', fontsize=11)
    ax.set_title(title, fontsize=12)
    ax.legend(fontsize=9)
    ax.grid(True, alpha=0.3)

plt.suptitle('Ridge vs Weighted Ridge RAPM Rankings',
             fontsize=13, fontweight='bold')
plt.tight_layout()
plt.show()

```



Results:

Variance Verification:

Multiplying margin by $\sqrt{n_pos}$ successfully stabilizes variance across possession buckets — the weighted standard deviation ranges narrowly from 113.30 to 132.64 compared to the original 15.78 to 91.55, confirming that $\text{Var}(\text{margin}) \propto 1/n_pos$ and that n_pos are the correct weights.

Weighted Ridge Results:

- The optimal alpha increased from 613.59 to 2,477.08 — indicating that weighting by possessions amplifies the signal in long stints, requiring stronger regularization.

- Top 20 median minutes remained high at 1,924 minutes, confirming that weighted Ridge continues to prioritize players with substantial playing time.

How Rankings Changed:

- Defensive players rose significantly — Draymond Green jumped from rank 7 to rank 2, and Giannis entered the top 10 at rank 7
- Offensive-focused players dropped — Trae Young (rank 3) and Coby White (rank 10) fell out of top 20 entirely. This suggests weighted Ridge better captures two-way players — those who contribute on both offense and defense across many possessions
- RAPM values compressed from a range of 3.17–6.12 to 2.50–4.54, reflecting more conservative and reliable estimates

Conclusion: Weighting by n_pos corrects the heteroscedasticity identified in Problem 2a, producing rankings that better reflect sustained performance across long stints rather than short high-margin bursts. By weighting each stint by its number of possessions — making short unreliable stints count less and long reliable stints count more in the RAPM calculation.

Problem 3: Interpreting Bootstrap Uncertainty

a. Calculate Confidence Intervals Using the Bootstrap: Suppose you are a general manager for a basketball team. You want to identify candidate players to add to your team, but you are unsure how certain to be about the ‘RAPM’ coefficients for the model. The bootstrap is a standard approach for calculating confidence intervals for model coefficients. Use either the ‘resample’ function from ‘sklearn’ or ‘np.random.choice’ along with a loop to calculate bootstrap samples of the model coefficients for each player. You may use the optimal ‘alpha’ found during the hyperparameter search in 2(b) for all bootstrap fits. Do not forget to resample the weights when bootstrapping!. For the top 20 players, display the RAPM estimates along with the confidence interval (calculate 92% intervals or some other high value that isn’t 95%). How much do the intervals of the top 20 overlap?

```

from sklearn.utils import resample

# 1. Bootstrap Setup
N_BOOTSTRAP = 100          # increase to 200-500 if computer is fast
ALPHA_FIXED = ridge_weighted.alpha_ # use optimal alpha from 2(b)
CONFIDENCE = 0.92          # 92% confidence interval
lower_pct = (1 - CONFIDENCE) / 2 # 0.04 = 4th percentile
upper_pct = 1 - (1 - CONFIDENCE) / 2 # 0.96 = 96th percentile

print(f"BOOTSTRAP CONFIGURATION")
print(f"{'='*45}")
print(f"  Samples:          {N_BOOTSTRAP}")
print(f"  Fixed alpha:       {ALPHA_FIXED:.2f}")
print(f"  Confidence level: {CONFIDENCE*100:.0f}%")
print(f"  Percentiles:      {lower_pct*100:.0f}th - {upper_pct*100:.0f}th")
print(f"  Data size:        {X.shape[0]:,} stints")

```

BOOTSTRAP CONFIGURATION

```

=====
Samples:          100
Fixed alpha:      2477.08
Confidence level: 92%
Percentiles:      4th - 96th
Data size:        32,358 stints

```

```

# 2. Run Bootstrap
bootstrap_coefs = np.zeros((N_BOOTSTRAP, X.shape[1]))

print(f"\nRunning {N_BOOTSTRAP} bootstrap samples...")

for i in range(N_BOOTSTRAP):
    # Resample X, y, AND weights together
    X_boot, y_boot, w_boot = resample(
        X, y, n_pos_weights,
        replace = True,
        random_state= i
    )

    # Fit Ridge with fixed optimal alpha
    with warnings.catch_warnings():

```

```

warnings.filterwarnings('ignore', category=RuntimeWarning)
ridge_boot = RidgeCV(
    alphas      = [ALPHA_FIXED],
    fit_intercept = True,
    cv          = 5
)
ridge_boot.fit(X_boot, y_boot, sample_weight=w_boot)

bootstrap_coefs[i, :] = ridge_boot.coef_

if (i + 1) % 10 == 0:
    print(f"  Completed {i+1}/{N_BOOTSTRAP} samples...")

print(f"Bootstrap complete!")
print(f"Coefficient matrix shape: {bootstrap_coefs.shape}")

```

Running 100 bootstrap samples...

```

Completed 10/100 samples...
Completed 20/100 samples...
Completed 30/100 samples...
Completed 40/100 samples...
Completed 50/100 samples...
Completed 60/100 samples...
Completed 70/100 samples...
Completed 80/100 samples...
Completed 90/100 samples...
Completed 100/100 samples...

```

Bootstrap complete!

Coefficient matrix shape: (100, 539)

```

# 3. Calculate Confidence Intervals
# Point estimates from weighted ridge
rapm_point = ridge_weighted.coef_

# Bootstrap CI
rapm_lower = np.percentile(bootstrap_coefs, lower_pct * 100, axis=0)
rapm_upper = np.percentile(bootstrap_coefs, upper_pct * 100, axis=0)
rapm_std    = bootstrap_coefs.std(axis=0)

# Build results dataframe

```



```

bootstrap_df = pd.DataFrame({
    'player_id': player_cols,
    'RAPM': rapm_point,
    'RAPM_lower': rapm_lower,
    'RAPM_upper': rapm_upper,
    'RAPM_std': rapm_std,
    'CI_width': rapm_upper - rapm_lower
})

# Merge with player names and minutes
bootstrap_df['player_id'] = bootstrap_df['player_id'].astype(str)
bootstrap_df = bootstrap_df.merge(player_names, on='player_id', how='left')
bootstrap_df['minutes_played'] = bootstrap_df['player_id'].map(minutes_played)

# Sort by RAPM
bootstrap_sorted = bootstrap_df.sort_values('RAPM', ascending=False).reset_index(drop=True)

# 4. Display Top 20 with Confidence Intervals
print(f"\nTOP 20 PLAYERS WITH {CONFIDENCE*100:.0f}% BOOTSTRAP CONFIDENCE INTERVALS")
print("=" * 85)
print(f"{'Rank':<5} {'Player':<22} {'RAPM':>7} "
      f"{'Lower':>8} {'Upper':>8} {'CI Width':>9} {'Minutes':>8}")
print("-" * 85)

top20_boot = bootstrap_sorted.head(20)
for i, row in top20_boot.iterrows():
    print(f"{i+1:<5} {str(row['player_name']):<22} "
          f"{row['RAPM']:>7.3f} "
          f"{row['RAPM_lower']:>8.3f} "
          f"{row['RAPM_upper']:>8.3f} "
          f"{row['CI_width']:>9.3f} "
          f"{row['minutes_played']:>8.0f}")

print(f"\n{'='*85}")
print(f"Average CI width (Top 20): {top20_boot['CI_width'].mean():.3f}")
print(f"Min CI width: {top20_boot['CI_width'].min():.3f}")
print(f"Max CI width: {top20_boot['CI_width'].max():.3f}")

```

TOP 20 PLAYERS WITH 92% BOOTSTRAP CONFIDENCE INTERVALS

=====

Rank	Player	RAPM	Lower	Upper	CI Width	Minutes
1	Nikola Jokic	4.541	2.536	6.494	3.958	2256
2	Draymond Green	4.463	2.727	6.696	3.969	2019
3	Joel Embiid	4.180	2.407	5.799	3.392	2249
4	Kawhi Leonard	3.903	2.066	5.665	3.599	1817
5	Anthony Davis	3.853	2.336	5.377	3.042	1881
6	Jrue Holiday	3.785	2.104	5.616	3.512	2097
7	Giannis Antetokounmpo	3.696	1.545	6.092	4.547	1841
8	Christian Koloko	3.538	1.106	5.308	4.202	758
9	Zion Williamson	3.403	1.499	5.211	3.712	859
10	Aaron Gordon	3.064	1.481	4.669	3.188	1967
11	Josh Hart	3.034	1.160	4.786	3.626	2475
12	Brook Lopez	2.843	0.336	4.413	4.076	2157
13	Jayson Tatum	2.652	0.705	4.782	4.077	2771
14	Desmond Bane	2.644	0.655	4.430	3.776	1706
15	Cameron Johnson	2.643	1.006	4.694	3.688	1073
16	Michael Porter Jr.	2.564	0.791	4.269	3.478	1686
17	Derrick White	2.560	0.264	4.342	4.078	2079
18	Nicolas Batum	2.543	0.805	4.068	3.263	1815
19	Franz Wagner	2.508	0.604	4.301	3.697	2293
20	Boban Marjanovic	2.498	1.043	4.092	3.049	167

```
=====
Average CI width (Top 20): 3.696
Min CI width:             3.042
Max CI width:             4.547
```

```
# 5. Visualize Confidence Intervals
fig, axes = plt.subplots(1, 2, figsize=(16, 8))

# Left: CI Plot for Top 20
top20_plot = bootstrap_sorted.head(20).copy()
top20_plot['short_name'] = top20_plot['player_name'].apply(
    lambda x: str(x).split()[-1]) # last name only

y_pos = np.arange(len(top20_plot))

axes[0].barh(y_pos,
             top20_plot['RAPM_upper'] - top20_plot['RAPM_lower'],
             left = top20_plot['RAPM_lower'],
             height= 0.6,
```

```

        color = 'steelblue', alpha=0.4, label='92% CI')

axes[0].scatter(top20_plot['RAPM'], y_pos,
                color='darkblue', s=50, zorder=5, label='RAPM estimate')


axes[0].axvline(x=0, color='red', linestyle='--', linewidth=1, alpha=0.7)
axes[0].set_yticks(y_pos)
axes[0].set_yticklabels(top20_plot['short_name'], fontsize=9)
axes[0].set_xlabel('RAPM (pts per 100 possessions)', fontsize=11)
axes[0].set_title(f'Top 20 RAPM with {CONFIDENCE*100:.0f}% Bootstrap CI\n'
                  f'(n={N_BOOTSTRAP} bootstrap samples)', fontsize=11)
axes[0].legend(fontsize=9)
axes[0].grid(True, alpha=0.3, axis='x')
axes[0].invert_yaxis()

# Right: CI Width vs Minutes
axes[1].scatter(bootstrap_df['minutes_played'],
                bootstrap_df['CI_width'],
                alpha=0.4, s=15, color='steelblue', label='All players')
axes[1].scatter(top20_plot['minutes_played'],
                top20_plot['CI_width'],
                color='red', s=50, zorder=5, label='Top 20')

axes[1].set_xlabel('Minutes Played', fontsize=11)
axes[1].set_ylabel('CI Width', fontsize=11)
axes[1].set_title('Confidence Interval Width vs Minutes Played\n'
                  '(More minutes = narrower CI)', fontsize=11)
axes[1].legend(fontsize=9)
axes[1].grid(True, alpha=0.3)

plt.suptitle('Bootstrap Uncertainty in RAPM Estimates',
             fontsize=13, fontweight='bold')
plt.tight_layout()
plt.show()

```



```
# 6. Overlap Analysis
print(f"\nCI OVERLAP ANALYSIS - TOP 20")
print("=" * 55)

# Check how many top 20 CIs overlap with rank 1 player
rank1_lower = top20_boot.iloc[0]['RAPM_lower']
rank1_upper = top20_boot.iloc[0]['RAPM_upper']
rank1_name = top20_boot.iloc[0]['player_name']

print(f"Rank 1: {rank1_name}")
print(f"CI: [{rank1_lower:.3f}, {rank1_upper:.3f}]")
print(f"\nOverlap with Rank 1:")
print("-" * 55)

overlap_count = 0
for i, row in top20_boot.iterrows():
    # Two intervals overlap if lower1 < upper2 AND upper1 > lower2
    overlaps = (row['RAPM_lower'] < rank1_upper and
                row['RAPM_upper'] > rank1_lower)
    if overlaps:
        overlap_count += 1
    rank = i + 1
    flag = " overlaps" if overlaps else " no overlap"
    print(f" Rank {rank:<3} {str(row['player_name']):<22} "
          f"[{row['RAPM_lower']:>6.3f}, {row['RAPM_upper']:>6.3f}] {flag}")
```

```
print(f"\n{overlap_count}/20 players overlap with Rank 1 CI")
```

CI OVERLAP ANALYSIS - TOP 20

```
=====
```

Rank 1: Nikola Jokic

CI: [2.536, 6.494]

Overlap with Rank 1:

```
-----
```

Rank 1	Nikola Jokic	[2.536, 6.494]	overlaps
Rank 2	Draymond Green	[2.727, 6.696]	overlaps
Rank 3	Joel Embiid	[2.407, 5.799]	overlaps
Rank 4	Kawhi Leonard	[2.066, 5.665]	overlaps
Rank 5	Anthony Davis	[2.336, 5.377]	overlaps
Rank 6	Jrue Holiday	[2.104, 5.616]	overlaps
Rank 7	Giannis Antetokounmpo	[1.545, 6.092]	overlaps
Rank 8	Christian Koloko	[1.106, 5.308]	overlaps
Rank 9	Zion Williamson	[1.499, 5.211]	overlaps
Rank 10	Aaron Gordon	[1.481, 4.669]	overlaps
Rank 11	Josh Hart	[1.160, 4.786]	overlaps
Rank 12	Brook Lopez	[0.336, 4.413]	overlaps
Rank 13	Jayson Tatum	[0.705, 4.782]	overlaps
Rank 14	Desmond Bane	[0.655, 4.430]	overlaps
Rank 15	Cameron Johnson	[1.006, 4.694]	overlaps
Rank 16	Michael Porter Jr.	[0.791, 4.269]	overlaps
Rank 17	Derrick White	[0.264, 4.342]	overlaps
Rank 18	Nicolas Batum	[0.805, 4.068]	overlaps
Rank 19	Franz Wagner	[0.604, 4.301]	overlaps
Rank 20	Boban Marjanovic	[1.043, 4.092]	overlaps

20/20 players overlap with Rank 1 CI

Bootstrap CI Results: The 92% bootstrap confidence intervals reveal substantial uncertainty in RAPM estimates across all top 20 players.

The average CI width is 3.696 points per 100 possessions — meaning most players' true RAPM could plausibly range across nearly 4 points in either direction.

Overlap Analysis:

All 20 top players' confidence intervals overlap with the rank 1 player (Jokic, CI: [2.536, 6.494]). This means we cannot statistically distinguish between any of the top 20 players with 92% confidence — the rankings reflect point estimates but carry significant uncertainty.

Key Observations:

- Players with highest lower bounds are safest picks: Anthony Davis (lower=2.336), Draymond Green (lower=2.727), and Jokic (lower=2.536) all have lower bounds above 2.0 — even in the worst case scenario they are well above average
- Christian Koloko (758 min) and Boban Marjanovic (167 min) appear in top 20 despite limited minutes — their wide CIs and low lower bounds make them high-risk picks -

All top 20 lower bounds are above zero, confirming every player is likely above average even accounting for uncertainty

Result: 20/20 players overlap with Rank 1

What it means:

We CANNOT statistically prove that Jokic (Rank 1) is better than any of the other top 20 players

All their confidence intervals share at least some common values with Jokic

As GM: Don't obsess over exact rankings Focus on lower bound instead

b. Impact of Minutes Played on Confidence Intervals: One potential use of a model like this is to find players who do not play much but who might do well with more opportunity. Calculate the standard errors of each coefficient from the bootstrap samples and make a scatterplot of standard error versus minutes played. Comment on the relationship between minutes played and standard error- does it make sense to you intuitively and/or statistically?

```
# Standard Error vs Minutes Plot
fig, ax = plt.subplots(figsize=(10, 6))

ax.scatter(bootstrap_df['minutes_played'],
           bootstrap_df['RAPM_std'],          # RAPM_std = standard error
           alpha=0.4, s=15, color='steelblue', label='All players')

ax.scatter(top20_boot['minutes_played'],
           top20_boot['RAPM_std'],
           color='red', s=50, zorder=5, label='Top 20')

# Add theoretical 1/sqrt(minutes) curve
min_range = np.linspace(
```

```

    bootstrap_df['minutes_played'].min() + 1,
    bootstrap_df['minutes_played'].max(), 300)

# Scale curve to data
scale = bootstrap_df['RAPM_std'].median() * np.sqrt(
    bootstrap_df['minutes_played'].median())
theoretical = scale / np.sqrt(min_range)

ax.plot(min_range, theoretical, color='red',
        linewidth=2, linestyle='--',
        label='Theoretical 1/√minutes')


ax.set_xlabel('Minutes Played', fontsize=12)
ax.set_ylabel('Standard Error of RAPM', fontsize=12)
ax.set_title('Standard Error vs Minutes Played\n'
            '(More minutes = more certain RAPM estimate)', fontsize=12)
ax.legend(fontsize=10)
ax.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()

# Print summary stats
print("STANDARD ERROR vs MINUTES SUMMARY")
print("=" * 50)

low_min = bootstrap_df[bootstrap_df['minutes_played'] < 200]['RAPM_std'].mean()
mid_min = bootstrap_df[(bootstrap_df['minutes_played'] >= 200) &
                        (bootstrap_df['minutes_played'] < 1000)]['RAPM_std'].mean()
high_min = bootstrap_df[bootstrap_df['minutes_played'] >= 1000]['RAPM_std'].mean()

print(f" Avg SE (<200 min):    {low_min:.4f} ← high uncertainty")
print(f" Avg SE (200-1000 min): {mid_min:.4f}")
print(f" Avg SE (>1000 min):    {high_min:.4f} ← low uncertainty")
print(f"\n Ratio high/low SE:    {low_min/high_min:.2f}x more uncertain")

```



STANDARD ERROR vs MINUTES SUMMARY

```
=====
Avg SE (<200 min):    0.6622  ← high uncertainty
Avg SE (200-1000 min):1.1309
Avg SE (>1000 min):   1.1033  ← low uncertainty

Ratio high/low SE:    0.60x more uncertain
```

This Seems Counterintuitive

Expected: more minutes = lower SE

Actual: very low minute players have LOWEST SE

Reason: Ridge regularization shrinks low-minute players coefficients toward zero

Players with <200 minutes: - Ridge shrinks their RAPM close to zero

- Every bootstrap sample also shrinks them to ~zero
- Coefficients barely move across samples
- Artificially LOW standard error
- This is NOT confidence — it is shrinkage

Players with 200-1000 minutes show the highest uncertainty as they have enough data to escape full shrinkage but not enough for stable estimates. Beyond 1000 minutes SE stabilizes, reflecting genuine statistical confidence from sufficient playing time.

c. Comparison to Bayesian Credible Intervals: Bayesian statistics (which you should have learned about briefly in DATA 606) is a field of statistics based on an interpretation of probability as referring to uncertainty about the real world rather than as the frequency of outcomes in repeated experiments. In the Bayesian framework, it is normal to talk about the probability distribution of model parameters, which leads to the concept of a ‘credible interval’, defined as an interval with a certain probability of containing the model parameter (a 92% credible interval would have a 92% chance of containing the model parameter). Credible intervals often correspond closely to confidence intervals captured using standard statistics, but this problem is one case where they diverge sharply. Ridge regression has an interpretation in Bayesian statistics, where the regularization parameter corresponds to the strength of a prior belief that the model coefficients have a normal distribution centered around zero and with variance inversely proportional to the regularization penalty. For ridge regression interpreted as Bayesian statistics, there is an exact formula for the standard errors and confidence intervals which we can use to contrast with the bootstrap estimate. I have provided code below to calculate the standard errors and the Bayesian credible intervals (if you are curious about the full details read this article). You will need to adapt this code to your variable names and data structures:

```
# 1. Prepare Variables
weights      = n_pos_weights      # possession weights
stint_mat    = X                  # player matrix
margin_vec   = y                  # target variable
alpha_opt    = ridge_weighted.alpha_opt # optimal alpha = 2477.08
model_ridge  = ridge_weighted     # fitted weighted ridge model

print(f"BAYESIAN CREDIBLE INTERVAL CALCULATION")
print(f"{'='*50}")
print(f"  Alpha (regularization): {alpha_opt:.2f}")
print(f"  Stint matrix shape:      {stint_mat.shape}")
print(f"  Weights shape:            {weights.shape}")
```

BAYESIAN CREDIBLE INTERVAL CALCULATION

```
=====
Alpha (regularization): 2477.08
Stint matrix shape:      (32358, 539)
Weights shape:           (32358,)
```

```

# 2. Calculate Bayesian Standard Errors
# Step 1: Weighted residual variance
var = np.average(
    (margin_vec - model_ridge.predict(stint_mat))**2,
    weights = weights
)

num_players = stint_mat.shape[1]
print(f" Residual variance:      {var:.4f}")
print(f" Number of players:      {num_players}")

# Step 2: A matrix = weighted data info + prior info
AMat = (stint_mat * weights.reshape(-1, 1)).T @ stint_mat + \
    alpha_opt * np.eye(num_players)

print(f" AMat shape:              {AMat.shape}")

# Step 3: Posterior covariance matrix
print(f"\n Inverting {num_players}x{num_players} matrix...")
posterior_covariance = var * np.linalg.inv(AMat)
print(f" Covariance matrix computed ")

# Step 4: Standard errors = sqrt of diagonal
se_bayesian = np.sqrt(np.diag(posterior_covariance))
print(f"\n Bayesian SE range: {se_bayesian.min():.4f} to {se_bayesian.max():.4f}")

```

```

Residual variance:      2054.7358
Number of players:      539
AMat shape:              (539, 539)

```

```

Inverting 539x539 matrix...
Covariance matrix computed

```

```

Bayesian SE range: 0.4827 to 0.9100

```

```

# 3. Add Bayesian SE to DataFrame
bootstrap_df['SE_bayesian'] = se_bayesian
bootstrap_df['SE_bootstrap'] = bootstrap_df['RAPM_std'] # already calculated
bootstrap_df['SE_ratio'] = (bootstrap_df['SE_bayesian'] /
    bootstrap_df['SE_bootstrap'])

```

```

# Bayesian credible intervals
z_92 = 1.75 # z-score for 92% interval
bootstrap_df['Bayes_lower'] = bootstrap_df['RAPM'] - z_92 * bootstrap_df['SE_bayesian']
bootstrap_df['Bayes_upper'] = bootstrap_df['RAPM'] + z_92 * bootstrap_df['SE_bayesian']

print(f"\nSE COMPARISON SUMMARY")
print(f"{'='*55}")
print(f"{'Metric':<30} {'Bayesian':>12} {'Bootstrap':>12}")
print(f"{'-'*55}")
print(f"{'Mean SE':<30} {se_bayesian.mean():>12.4f} "
      f"{bootstrap_df['SE_bootstrap'].mean():>12.4f}")
print(f"{'Median SE':<30} {np.median(se_bayesian):>12.4f} "
      f"{bootstrap_df['SE_bootstrap'].median():>12.4f}")
print(f"{'Min SE':<30} {se_bayesian.min():>12.4f} "
      f"{bootstrap_df['SE_bootstrap'].min():>12.4f}")
print(f"{'Max SE':<30} {se_bayesian.max():>12.4f} "
      f"{bootstrap_df['SE_bootstrap'].max():>12.4f}")

```

SE COMPARISON SUMMARY

```

=====
Metric                                Bayesian    Bootstrap
-----
Mean SE                             0.6581      1.0235
Median SE                           0.6157      1.0801
Min SE                              0.4827      0.0872
Max SE                              0.9100      1.3802

```

```

# 4. Plot Bayesian SE vs Minutes
fig, axes = plt.subplots(1, 3, figsize=(18, 6))

# Left: Bayesian SE vs Minutes
axes[0].scatter(bootstrap_df['minutes_played'],
                bootstrap_df['SE_bayesian'],
                alpha=0.4, s=15, color='steelblue')
axes[0].set_xlabel('Minutes Played', fontsize=11)
axes[0].set_ylabel('Standard Error', fontsize=11)
axes[0].set_title('Bayesian SE vs Minutes', fontsize=12)
axes[0].grid(True, alpha=0.3)

# Middle: Bootstrap SE vs Minutes

```


```

axes[1].scatter(bootstrap_df['minutes_played'],
                bootstrap_df['SE_bootstrap'],
                alpha=0.4, s=15, color='darkorange')
axes[1].set_xlabel('Minutes Played', fontsize=11)
axes[1].set_ylabel('Standard Error', fontsize=11)
axes[1].set_title('Bootstrap SE vs Minutes', fontsize=12)
axes[1].grid(True, alpha=0.3)

# Right: Both overlaid for comparison
axes[2].scatter(bootstrap_df['minutes_played'],
                bootstrap_df['SE_bayesian'],
                alpha=0.4, s=15, color='steelblue',
                label='Bayesian')
axes[2].scatter(bootstrap_df['minutes_played'],
                bootstrap_df['SE_bootstrap'],
                alpha=0.4, s=15, color='darkorange',
                label='Bootstrap')
axes[2].set_xlabel('Minutes Played', fontsize=11)
axes[2].set_ylabel('Standard Error', fontsize=11)
axes[2].set_title('Bayesian vs Bootstrap SE\n(Comparison)', fontsize=12)
axes[2].legend(fontsize=10)
axes[2].grid(True, alpha=0.3)

plt.suptitle('Bayesian vs Bootstrap Standard Errors',
             fontsize=13, fontweight='bold')
plt.tight_layout()
plt.show()

```



```

# 5. Agreement Analysis
print(f"\nAGREEMENT ANALYSIS BY MINUTES BUCKET")
print(f"{'='*65}")
print(f"{'Minutes':<15} {'Bayes SE':>10} {'Boot SE':>10} "
      f"{'Ratio':>8} {'Agreement':>12}")
print(f"{'-'*65}")

buckets = [
    ('<200', 0, 200),
    ('200-500', 200, 500),
    ('500-1000', 500, 1000),
    ('>1000', 1000, float('inf'))
]

for label, low, high in buckets:
    mask = ((bootstrap_df['minutes_played'] >= low) &
            (bootstrap_df['minutes_played'] < high))
    subset = bootstrap_df[mask]

    if len(subset) == 0:
        continue

    bayes_se = subset['SE_bayesian'].mean()
    boot_se = subset['SE_bootstrap'].mean()
    ratio = bayes_se / boot_se

    if 0.8 <= ratio <= 1.2:
        agreement = "HIGH"
    elif 0.5 <= ratio <= 1.5:
        agreement = "MODERATE"
    else:
        agreement = "LOW"

    print(f"{label:<15} {bayes_se:>10.4f} {boot_se:>10.4f} "
          f"{ratio:>8.2f} {agreement:>12}")

```

AGREEMENT ANALYSIS BY MINUTES BUCKET

```

=====
Minutes          Bayes SE    Boot SE    Ratio    Agreement
-----
<200              0.8599     0.6622     1.30     MODERATE

```

200-500	0.7439	1.0825	0.69	MODERATE
500-1000	0.6428	1.1595	0.55	MODERATE
>1000	0.5614	1.1033	0.51	MODERATE

Overall SE Comparison

- Bayesian SE range: 0.48 to 0.91 $\leftarrow$ narrow, consistent
- Bootstrap SE range: 0.09 to 1.38 $\leftarrow$ wide, inconsistent
- Bayesian mean SE: 0.66
- Bootstrap mean SE: 1.02 $\leftarrow$ 55% larger on average

Minutes Bayes SE Boot SE Ratio

- <200 0.86 0.66 1.30 $\leftarrow$ Bayes HIGHER
- 200-500 0.74 1.08 0.69 $\leftarrow$ Bootstrap HIGHER
- 500-1000 0.64 1.16 0.55 $\leftarrow$ Bootstrap HIGHER
- >1000 0.56 1.10 0.51 $\leftarrow$ Bootstrap HIGHER


```
fig, ax = plt.subplots(figsize=(10, 6))

ax.scatter(bootstrap_df['minutes_played'],
           bootstrap_df['SE_bayesian'],
           alpha=0.4, s=15, color='steelblue',
           label='Bayesian SE')
ax.scatter(bootstrap_df['minutes_played'],
           bootstrap_df['SE_bootstrap'],
           alpha=0.4, s=15, color='darkorange',
           label='Bootstrap SE')

# Mark the crossover point
ax.axvline(x=200, color='red', linestyle='--',
           linewidth=1.5, label='Crossover (~200 min)')

ax.set_xlabel('Minutes Played', fontsize=12)
ax.set_ylabel('Standard Error', fontsize=12)
ax.set_title('Bayesian vs Bootstrap SE\nCrossover at ~200 minutes',
            fontsize=12)
ax.legend(fontsize=10)
ax.grid(True, alpha=0.3)
```

```
plt.tight_layout()
plt.show()
```



Crossover happened at 200 minutes

Agreement:

Highest agreement: <200 minutes (ratio=1.30)

- Both methods closest here
- Still only MODERATE agreement
- No region shows truly HIGH agreement

Disagreement:

Highest disagreement: >1000 minutes (ratio=0.51) → Bootstrap SE twice as large as Bayesian SE

- Completely opposite of what was expected

Which is More Realistic:

For low minute players (<200 min):

Bayesian SE: 0.86 ← higher, more uncertain
Bootstrap SE: 0.66 ← lower, falsely confident

Bayesian is more realistic here

- Correctly recognizes limited data = high uncertainty
- Bootstrap is fooled by Ridge shrinkage
- Low minute players shrunk to zero consistently
- Gives artificially low SE

For high minute players (>1000 min):

Bayesian SE: 0.56 ← lower, more confident
Bootstrap SE: 1.10 ← higher, more uncertain

- Bayesian is more realistic here too
- 1000+ minutes = lots of data = should be confident
 - Bootstrap overestimates uncertainty for star players
 - Resampling creates variation that isn't really there
 - Bayesian formula correctly uses all information

Conclusion:

Bayesian credible intervals are more realistic
across ALL minute ranges because:

1. Low minute players:
 - Bayesian correctly shows HIGH uncertainty
 - Bootstrap falsely shows low SE due to shrinkage
2. High minute players:
 - Bayesian correctly shows LOW uncertainty
 - Bootstrap overestimates uncertainty

The Bayesian formula accounts for the exact
information content in the data mathematically,
while Bootstrap is distorted by Ridge
regularization at both extremes of minutes played

As a GM: Trust Bayesian intervals over Bootstrap
for both identifying reliable stars AND
flagging uncertain low-minute players

Extra Credit (5 Points): Ridge versus Lasso and Confidence/Credible Interval Calculations

a. An alternative to ridge regression that is used for variable selection is called the Lasso. The Lasso differs penalty causes a large number of model coefficients to be zero, making it a good tool for variable selection and for creating interpretable models. Use 'LassoCV' to calculate the 'RAPM' coefficients. I recommend using regularization weights that range between and (the scale needs to be different compared to ridge regression). Compare the 'RAPM' values calculated with lasso to those calculated with ridge regression. Are there any notable differences in the top 20 players? Find the 10 players with the largest difference between lasso RAPM and Ridge RAPM in both directions (ridge greater than lasso and lasso greater than ridge). Within the players where there was large disagreement, consider how the players performed in the next NBA season (you may do this however you like, I recommend reading media reports about those players during the next season) and determine which model was more correct when it comes to players with large disagreements.

Intuition

****The Key Difference****

RIDGE:	LASSO:
Shrinks ALL coefficients toward zero	Sets MANY coefficients exactly TO zero
Like turning down a dimmer switch	Like turning OFF a light switch
Player with little data: Ridge RAPM $\rightarrow$ 0.5 (small but exists)	Player with little data: Lasso RAPM $\rightarrow$ 0.0 exactly (completely removed)

Result: Lasso keeps only players with enough data to make a judgment = automatic variable selection

```
from sklearn.linear_model import LassoCV
import warnings

# 1. Define Alpha Range for Lasso
# Lasso needs different scale than Ridge
lasso_alphas = np.logspace(-3, 1, 100)

print(f"LASSO CONFIGURATION")
print(f"{'='*45}")
```

```
print(f" Alpha range: {lasso_alphas[0]:.4f} to {lasso_alphas[-1]:.2f}")
print(f" Number of alphas: {len(lasso_alphas)}")
```

LASSO CONFIGURATION

```
=====
```

```
Alpha range: 0.0010 to 10.00
Number of alphas: 100
```

```
# 2. Fit Lasso with Cross Validation
with warnings.catch_warnings():
    warnings.filterwarnings('ignore')
    lasso = LassoCV(
        alphas          = lasso_alphas,
        fit_intercept    = True,
        cv               = 5,
        max_iter         = 10000      # Lasso needs more iterations
    )
    lasso.fit(X, y, sample_weight=n_pos_weights)

print(f"LASSO CV RESULTS")
print(f"{'='*50}")
print(f" Optimal alpha:          {lasso.alpha_:.6f}")
print(f" Log10(alpha):            {np.log10(lasso.alpha_):.2f}")
print(f" Intercept (HCA):         {lasso.intercept_:.4f}")
print(f" Non-zero coefs:          {(lasso.coef_ != 0).sum()}")
print(f" Zero coefs:              {(lasso.coef_ == 0).sum()}")
print(f" % players zeroed out:    {(lasso.coef_ == 0).sum()/len(lasso.coef_)*100:.1f}%")
```

LASSO CV RESULTS

```
=====
```

```
Optimal alpha:          0.025950
Log10(alpha):           -1.59
Intercept (HCA):        1.0779
Non-zero coefs:         197
Zero coefs:             342
% players zeroed out:   63.5%
```

```
# 3. Build Lasso RAPM DataFrame
lasso_rapm_df = pd.DataFrame({
    'player_id': player_cols,
```

```

    'RAPM_lasso': lasso.coef_
})

lasso_rapm_df['player_id'] = lasso_rapm_df['player_id'].astype(str)
lasso_rapm_df = lasso_rapm_df.merge(player_names, on='player_id', how='left')
lasso_rapm_df['minutes_played'] = lasso_rapm_df['player_id'].map(minutes_played)

lasso_rapm_sorted = lasso_rapm_df.sort_values(
    'RAPM_lasso', ascending=False).reset_index(drop=True)

print(f"\nTOP 20 PLAYERS BY LASSO RAPM")
print(f"{'='*65}")
print(f"{'Rank':<5} {'Player':<25} {'RAPM_Lasso':>12} {'Minutes':>10}")
print(f"{'-'*65}")
for i, row in lasso_rapm_sorted.head(20).iterrows():
    print(f"{i+1:<5} {str(row['player_name']):<25} "
          f"{row['RAPM_lasso']:>12.4f} "
          f"{row['minutes_played']:>10.0f}")

```

TOP 20 PLAYERS BY LASSO RAPM

Rank	Player	RAPM_Lasso	Minutes
1	Draymond Green	5.8801	2019
2	Joel Embiid	5.8795	2249
3	Nikola Jokic	5.5918	2256
4	Kawhi Leonard	4.4256	1817
5	Anthony Davis	4.2591	1881
6	Christian Koloko	4.1152	758
7	Jrue Holiday	3.9740	2097
8	Zion Williamson	3.8889	859
9	Giannis Antetokounmpo	3.5087	1841
10	Cameron Johnson	3.3588	1073
11	Josh Hart	3.1860	2475
12	Derrick White	3.0914	2079
13	Julius Randle	3.0298	2957
14	Jayson Tatum	2.9368	2771
15	Shai Gilgeous-Alexander	2.6284	2620
16	Lauri Markkanen	2.6113	2130
17	Desmond Bane	2.6057	1706
18	Darius Garland	2.5780	2210

19	Aaron Gordon	2.5567	1967
20	Kevin Durant	2.5113	1729

```
# 4. Compare Ridge vs Lasso Top 20
ridge_top20 = weighted_rapm_sorted.head(20).reset_index(drop=True)
lasso_top20 = lasso_rapm_sorted.head(20).reset_index(drop=True)

print(f"\nRIDGE vs LASSO TOP 20 COMPARISON")
print(f"{'='*90}")
print(f"{'--- RIDGE ---':<45} {'--- LASSO ---':<45}")
print(f"{'Rank':<5} {'Player':<22} {'RAPM':>8} {'Min':>6}  "
      f"{'Rank':<5} {'Player':<22} {'RAPM':>8} {'Min':>6}")
print(f"{'='*90}")

for i in range(20):
    r = ridge_top20.iloc[i]
    l = lasso_top20.iloc[i]
    print(f"{i+1:<5} {str(r['player_name']):<22} "
          f"{r['RAPM']:>8.3f} {r['minutes_played']:>6.0f}  "
          f"{i+1:<5} {str(l['player_name']):<22} "
          f"{l['RAPM_lasso']:>8.3f} {l['minutes_played']:>6.0f}")
```

RIDGE vs LASSO TOP 20 COMPARISON

--- RIDGE ---				--- LASSO ---			
Rank	Player	RAPM	Min	Rank	Player	RAPM	Min
1	Nikola Jokic	4.541	2256	1	Draymond Green	5.880	2019
2	Draymond Green	4.463	2019	2	Joel Embiid	5.880	2249
3	Joel Embiid	4.180	2249	3	Nikola Jokic	5.592	2256
4	Kawhi Leonard	3.903	1817	4	Kawhi Leonard	4.426	1817
5	Anthony Davis	3.853	1881	5	Anthony Davis	4.259	1881
6	Jrue Holiday	3.785	2097	6	Christian Koloko	4.115	758
7	Giannis Antetokounmpo	3.696	1841	7	Jrue Holiday	3.974	2097
8	Christian Koloko	3.538	758	8	Zion Williamson	3.889	859
9	Zion Williamson	3.403	859	9	Giannis Antetokounmpo	3.509	1841
10	Aaron Gordon	3.064	1967	10	Cameron Johnson	3.359	1073
11	Josh Hart	3.034	2475	11	Josh Hart	3.186	2475
12	Brook Lopez	2.843	2157	12	Derrick White	3.091	2079
13	Jayson Tatum	2.652	2771	13	Julius Randle	3.030	2957
14	Desmond Bane	2.644	1706	14	Jayson Tatum	2.937	2771

15	Cameron Johnson	2.643	1073	15	Shai Gilgeous-Alexander	2.628	2620
16	Michael Porter Jr.	2.564	1686	16	Lauri Markkanen	2.611	2130
17	Derrick White	2.560	2079	17	Desmond Bane	2.606	1706
18	Nicolas Batum	2.543	1815	18	Darius Garland	2.578	2210
19	Franz Wagner	2.508	2293	19	Aaron Gordon	2.557	1967
20	Boban Marjanovic	2.498	167	20	Kevin Durant	2.511	1729

Top 20 Comparison Analysis

Where They Agree

Both have in top 20:

→ Jokic, Draymond, Embiid, Kawhi

→ Anthony Davis, Jrue Holiday

→ Giannis, Zion, Josh Hart

→ Jayson Tatum, Derrick White

Core elite players identified by BOTH models = high confidence these are truly good players

```
# 5. Find Largest Disagreements
# Merge Ridge and Lasso on player_id
comparison_df = weighted_rapm_df[['player_id', 'player_name',
                                   'RAPM', 'minutes_played']].copy()

comparison_df = comparison_df.merge(
    lasso_rapm_df[['player_id', 'RAPM_lasso']],
    on='player_id', how='inner'
)

comparison_df['difference'] = comparison_df['RAPM'] - comparison_df['RAPM_lasso']

# Ridge > Lasso (top 10)
ridge_greater = comparison_df.nlargest(10, 'difference')

# Lasso > Ridge (top 10)
lasso_greater = comparison_df.nsmallest(10, 'difference')

print(f"\nTOP 10: RIDGE RAPM > LASSO RAPM")
print(f"{'='*75}")
print(f"{'Player':<22} {'Ridge':>8} {'Lasso':>8} {'Diff':>8} {'Minutes':>8}")
print(f"{'-'*75}")
for _, row in ridge_greater.iterrows():
    print(f"{str(row['player_name']):<22} "
```

```

f"{row['RAPM']:>8.3f} "
f"{row['RAPM_lasso']:>8.3f} "
f"{row['difference']:>+8.3f} "
f"{row['minutes_played']:>8.0f}")

print(f"\nTOP 10: LASSO RAPM > RIDGE RAPM")
print(f"{'='*75}")
print(f"{'Player':<22} {'Ridge':>8} {'Lasso':>8} {'Diff':>8} {'Minutes':>8}")
print(f"{'-'*75}")
for _, row in lasso_greater.iterrows():
    print(f"{str(row['player_name']):<22} "
          f"{row['RAPM']:>8.3f} "
          f"{row['RAPM_lasso']:>8.3f} "
          f"{row['difference']:>+8.3f} "
          f"{row['minutes_played']:>8.0f}")

```

TOP 10: RIDGE RAPM > LASSO RAPM

```

=====
Player                                Ridge    Lasso    Diff    Minutes
-----
Kira Lewis Jr.                        1.844    0.000    +1.844    230
Jared Rhoden                          1.827    0.000    +1.827    184
McKinley Wright IV                   -3.057   -4.774    +1.716    335
Charles Bassey                       1.671    0.000    +1.671    556
Clint Capela                         1.525    0.000    +1.525   1696
Rodney McGruder                      1.501    0.000    +1.501    512
Moses Brown                          1.492    0.000    +1.492    275
Gary Payton II                       1.463    0.000    +1.463    380
Frank Ntilikina                      1.477    0.035    +1.442    618
Gordon Hayward                       1.382    0.000    +1.382   1450

```

TOP 10: LASSO RAPM > RIDGE RAPM

```

=====
Player                                Ridge    Lasso    Diff    Minutes
-----
Johnny Juzang                         -1.932   -0.051   -1.881    222
Jeff Dowtin Jr.                      -1.854   -0.007   -1.847    238
Matthew Dellavedova                  -1.741   -0.000   -1.741    219
Ryan Rollins                         -1.717   -0.000   -1.717     65
Joel Embiid                          4.180    5.880   -1.699   2249
Saben Lee                           -1.678   -0.000   -1.678    342

```

Payton Pritchard	-1.713	-0.037	-1.676	618
Davis Bertans	-1.669	-0.000	-1.669	422
Udoka Azubuike	-1.644	-0.000	-1.644	355
JaVale McGee	-1.594	-0.000	-1.594	335

Disagreement Analysis

Ridge > Lasso (Ridge gives more credit)

8 out of 10 players have Lasso RAPM = 0 These are players Ridge credits but Lasso completely ignores:

- Clint Capela: Ridge=1.525 Lasso=0 1,696 min
- Gordon Hayward: Ridge=1.382 Lasso=0 1,450 min

Lasso > Ridge (Lasso gives more credit)

Interesting pattern - mostly NEGATIVE players:

- Johnny Juzang: Ridge=-1.932 Lasso=-0.051
- Jeff Downtin Jr.: Ridge=-1.854 Lasso=-0.007

Lasso zeroes these OUT (sets to ~0) Ridge says they are ACTIVELY HURTING team

Joel Embiid exception:

Ridge=4.180 Lasso=5.880 Diff=-1.699

Lasso gives Embiid MORE credit than Ridge

- Lasso sees his contribution as more unique
- Less shared with teammates

```
# 6. Scatter Plot Ridge vs Lasso RAPM
fig, axes = plt.subplots(1, 2, figsize=(14, 6))

# Left: Ridge vs Lasso scatter
axes[0].scatter(comparison_df['RAPM'],
                comparison_df['RAPM_lasso'],
                alpha=0.4, s=15, color='steelblue')
axes[0].plot([comparison_df['RAPM'].min(),
              comparison_df['RAPM'].max()],
             [comparison_df['RAPM'].min(),
              comparison_df['RAPM'].max()],
             'r--', linewidth=1.5, label='Perfect agreement')
axes[0].set_xlabel('Ridge RAPM', fontsize=11)
```


```

axes[0].set_ylabel('Lasso RAPM', fontsize=11)
axes[0].set_title('Ridge vs Lasso RAPM\n(Points on line = perfect agreement)',
                  fontsize=11)
axes[0].legend(fontsize=9)
axes[0].grid(True, alpha=0.3)

# Right: Difference vs Minutes
axes[1].scatter(comparison_df['minutes_played'],
                comparison_df['difference'],
                alpha=0.4, s=15, color='steelblue')
axes[1].axhline(y=0, color='red', linestyle='--', linewidth=1.5)
axes[1].set_xlabel('Minutes Played', fontsize=11)
axes[1].set_ylabel('Ridge RAPM - Lasso RAPM', fontsize=11)
axes[1].set_title('Disagreement vs Minutes Played\n'
                  '(Above zero = Ridge higher)',
                  fontsize=11)
axes[1].grid(True, alpha=0.3)

plt.suptitle('Ridge vs Lasso RAPM Comparison',
             fontsize=13, fontweight='bold')
plt.tight_layout()
plt.show()

```



```

# 7. Summary Statistics
print(f"\nSUMMARY: RIDGE vs LASSO")
print(f"{'='*55}")
print(f"{'Metric':<35} {'Ridge':>10} {'Lasso':>10}")

```



```

print(f"{'-'*55}")
print(f"{'Non-zero coefficients':<35} "
      f"{'{(weighted_rapm_df['RAPM'] != 0).sum():>10} "
      f"{'{(lasso.coef_ != 0).sum():>10}"}")
print(f"{'Mean RAPM (all players)':<35} "
      f"{'{(weighted_rapm_df['RAPM']).mean():>10.4f} "
      f"{'{(lasso_rapm_df['RAPM_lasso']).mean():>10.4f}"}")
print(f"{'Max RAPM':<35} "
      f"{'{(weighted_rapm_df['RAPM']).max():>10.4f} "
      f"{'{(lasso_rapm_df['RAPM_lasso']).max():>10.4f}"}")
print(f"{'Min RAPM':<35} "
      f"{'{(weighted_rapm_df['RAPM']).min():>10.4f} "
      f"{'{(lasso_rapm_df['RAPM_lasso']).min():>10.4f}"}")
print(f"{'Optimal Alpha':<35} "
      f"{'{(ridge_weighted.alpha_>10.2f} "
      f"{'{(lasso.alpha_>10.6f}"}")
print(f"{'Players zeroed out':<35} "
      f"{'{'0':>10} "
      f"{'{(lasso.coef_ == 0).sum():>10}"}")

```

SUMMARY: RIDGE vs LASSO

Metric	Ridge	Lasso
Non-zero coefficients	539	197
Mean RAPM (all players)	-0.0001	0.0772
Max RAPM	4.5407	5.8801
Min RAPM	-3.8503	-4.9552
Optimal Alpha	2477.08	0.025950
Players zeroed out	0	342

Summary Findings 1. Aggressive Variable Selection: Lasso zeroed out 342 of 539 players (63.5%), keeping only 197 players with non-zero coefficients. This confirms that the majority of NBA players in this dataset do not have enough unique statistical signal to distinguish their individual contribution from noise — Ridge assigned all 539 players small coefficients while Lasso made a definitive judgment that most bench and low-minute players cannot be reliably evaluated.

2. More Extreme RAPM Range: Lasso produced more extreme values in both directions — max RAPM of 5.88 vs Ridge's 4.54, and min RAPM of -4.96 vs Ridge's -3.85. By zeroing

out 342 players, Lasso concentrates all the explanatory signal into the remaining 197 players, giving them larger coefficients in both positive and negative directions.

3. Optimal Alpha Scale Difference: The optimal alpha for Lasso (0.026) is dramatically smaller than Ridge (2477.08) — nearly 100,000 times smaller. This reflects a fundamental difference in how the two penalties work: Lasso's absolute value penalty is inherently stronger at producing zero coefficients, so it requires far less regularization strength to achieve the same level of shrinkage that Ridge needs a very large alpha to accomplish.

Conclusion Lasso is better at:

- Variable selection (removing noise players)
- Identifying independently elite players
- Surfacing stars hidden in Ridge top 20 (Shai, Markkanen, Durant)

Ridge is better at:

- Capturing role players and team contributors
- Distributing credit among teammates
- Not over-penalizing low minute players who may genuinely contribute

For a GM: Use BOTH models together Players in BOTH top 20 = highest confidence Players in only one = worth investigating further

Online search and results:

Player name: CLINT CAPELA

Player 2023-2024 stats link: https://www.espn.com/nba/player/stats/_/id/3102529/clint-capela

Stats Regular Season Averages

GP = 73 ← MOST IMPORTANT

MIN = 25.8 ← SECOND MOST IMPORTANT

PTS = 11.5 ← SUPPORTING EVIDENCE

Conclusion based on stats Clint Capela

played 73 games averaging 25.8 minutes as a starter in 2023-24,

confirming Ridge was more correct — his contribution was real and Lasso incorrectly zeroed him out.

Player Name: Gordon Hayward (Ridge=1.382, Lasso=0)

player 2023-2024 stats link: https://www.espn.com/nba/player/stats/_/id/4249/gordon-hayward

Charlotte (CHA): 25 GP 25 GS 31.9 MIN

Oklahoma City(OKC): 26 GP 3 GS 17.2 MIN

Conclusion based on stats

Gordon Hayward started 25 games for Charlotte but after being traded to OKC his role diminished sharply to just 17.2 minutes off the bench, suggesting *Lasso* was more correct to discount his value.

Player name: Gary Payton II

Analysis (Ridge = 1.463, Lasso = 0)

player (2023-2024) stats link: https://www.espn.com/nba/player/stats/_/id/3134903/gary-payton-ii 2023-24 GS: 44 GP 0 GS 15.5 MIN 2024-25 GS: 62 GP 11 GS 15.0 MIN

Conclusion based on stats

Gary Payton II played a limited bench role in 2023-24 (44 games, 0 starts) but grew into a larger role in 2024-25,

suggesting Ridge detected real but developing value that Lasso missed.

His continued NBA presence across two more seasons confirms he was a real contributor — not just noise that should be zeroed out.

Defensive specialists like Payton are exactly the type of player Ridge captures better than Lasso.

b. There is a subtle conceptual flaw in how the confidence or credible intervals were calculated in this lab. This is apparent if you paid very close attention during the meetup. Can you tell me what it is?

What we did

Step 1: Used RidgeCV to find optimal alpha $\rightarrow$ alpha = 2477.08

Step 2: Used that SAME alpha for ALL bootstrap samples

Step 3: Calculated confidence intervals from bootstrap coefficients

Flaw

We fixed alpha = 2477.08 for every bootstrap sample

BUT alpha should be RE-SELECTED for each bootstrap sample separately!

Each bootstrap sample is a different “version” of the data

$\rightarrow$ should get its own optimal alpha via cross validation

Why this Matters

WHAT WE DID:

Bootstrap sample 1 → fixed alpha 2477.08

Bootstrap sample 2 → fixed alpha 2477.08

Bootstrap sample 3 → fixed alpha 2477.08 ...

→ Underestimates uncertainty in alpha itself

WHAT WE SHOULD DO:

Bootstrap sample 1 → RidgeCV → alpha_1

Bootstrap sample 2 → RidgeCV → alpha_2

Bootstrap sample 3 → RidgeCV → alpha_3 ...

→ Captures FULL uncertainty including uncertainty in hyperparameter selection

Consequence

By fixing alpha we assumed: “We are 100% certain 2477.08 is correct”

But alpha itself has uncertainty! Different data samples would select different optimal alphas

Result:

Our confidence intervals are TOO NARROW — they underestimate the true uncertainty in RAPM estimates
